

# Machine Learning for Graphics and Vision

Prof. Dr.-Ing. Andreas Geiger

Autonomous Vision Group  
MPI-IS / University of Tübingen

April 19, 2018



University of Tübingen  
MPI for Intelligent Systems  

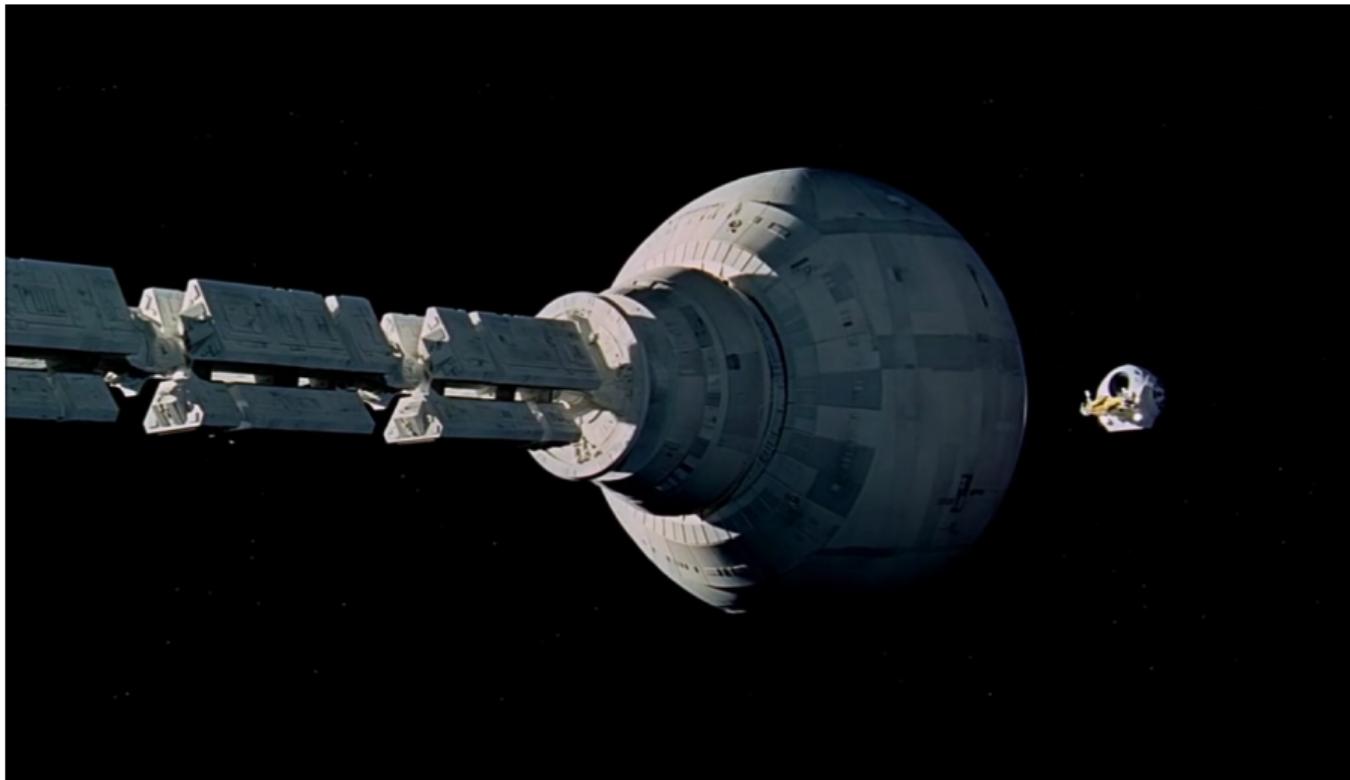
---

**Autonomous Vision Group**



# What is Artificial Intelligence?

# Artificial Intelligence

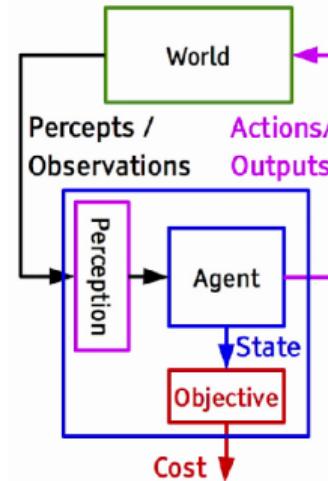


# Artificial Intelligence

"An attempt will be made to find how to make machines use language, form abstractions and concepts, solve kinds of problems now reserved for humans, and improve themselves."

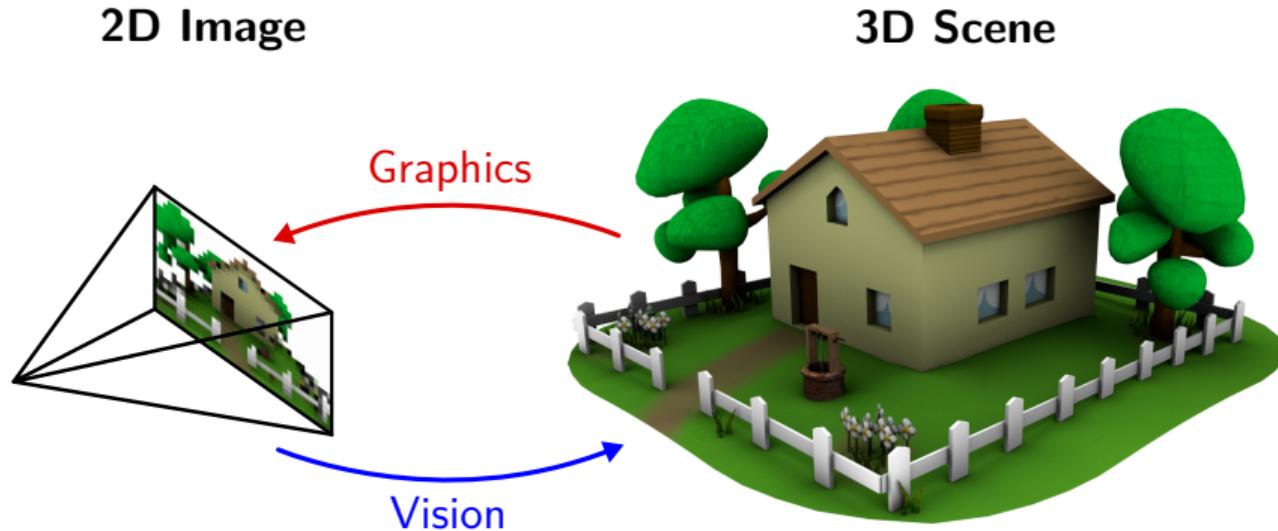
[John McCarthy]

- ▶ Machine Learning
- ▶ Computer Vision
- ▶ Natural Language Processing
- ▶ Robotics & Control
- ▶ Graphics
- ▶ Art, Industry 4.0, Education ...



But: No unique definition (e.g., OCR)

# Computer Graphics vs. Computer Vision

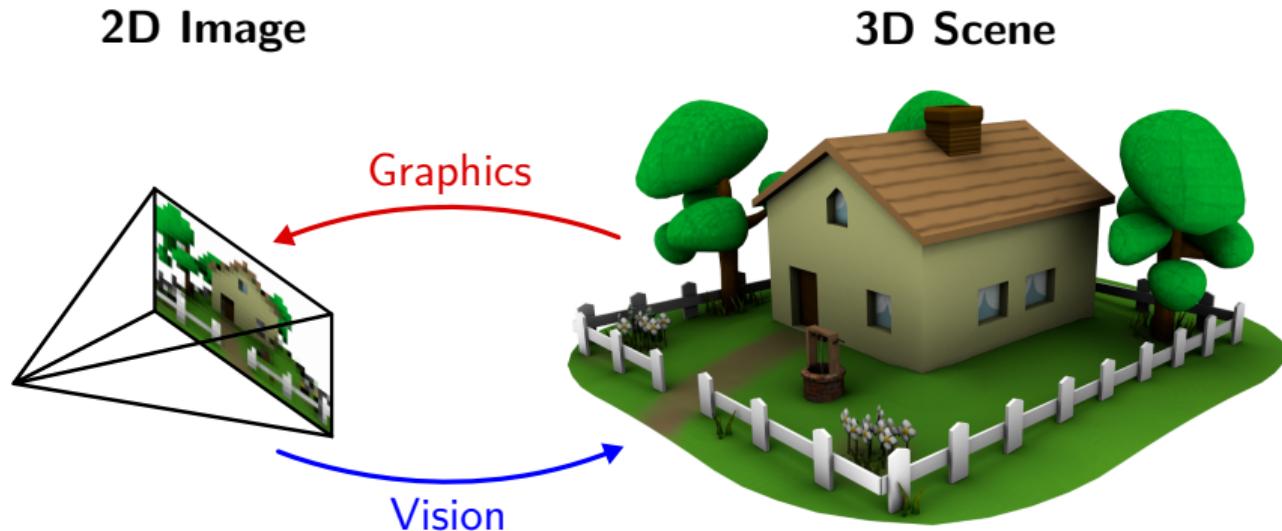


Pixel Matrix

217	191	252	255	239
102	80	200	146	138
159	94	91	121	138
179	106	136	85	41
115	129	83	112	67
94	114	105	111	89

Objects	Material
Shape/Geometry	Motion
Semantics	3D Pose

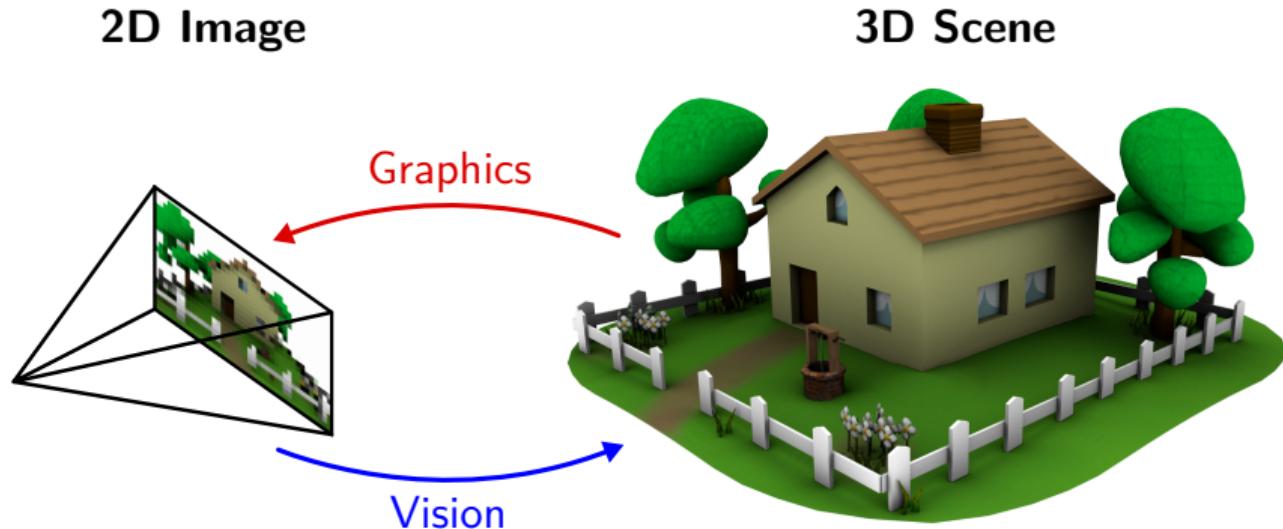
# Computer Graphics vs. Computer Vision



Computer Vision is an ill-posed inverse problem:

- ▶ Many 3D scenes yield the same 2D image
- ▶ Additional constraints (knowledge about world) required

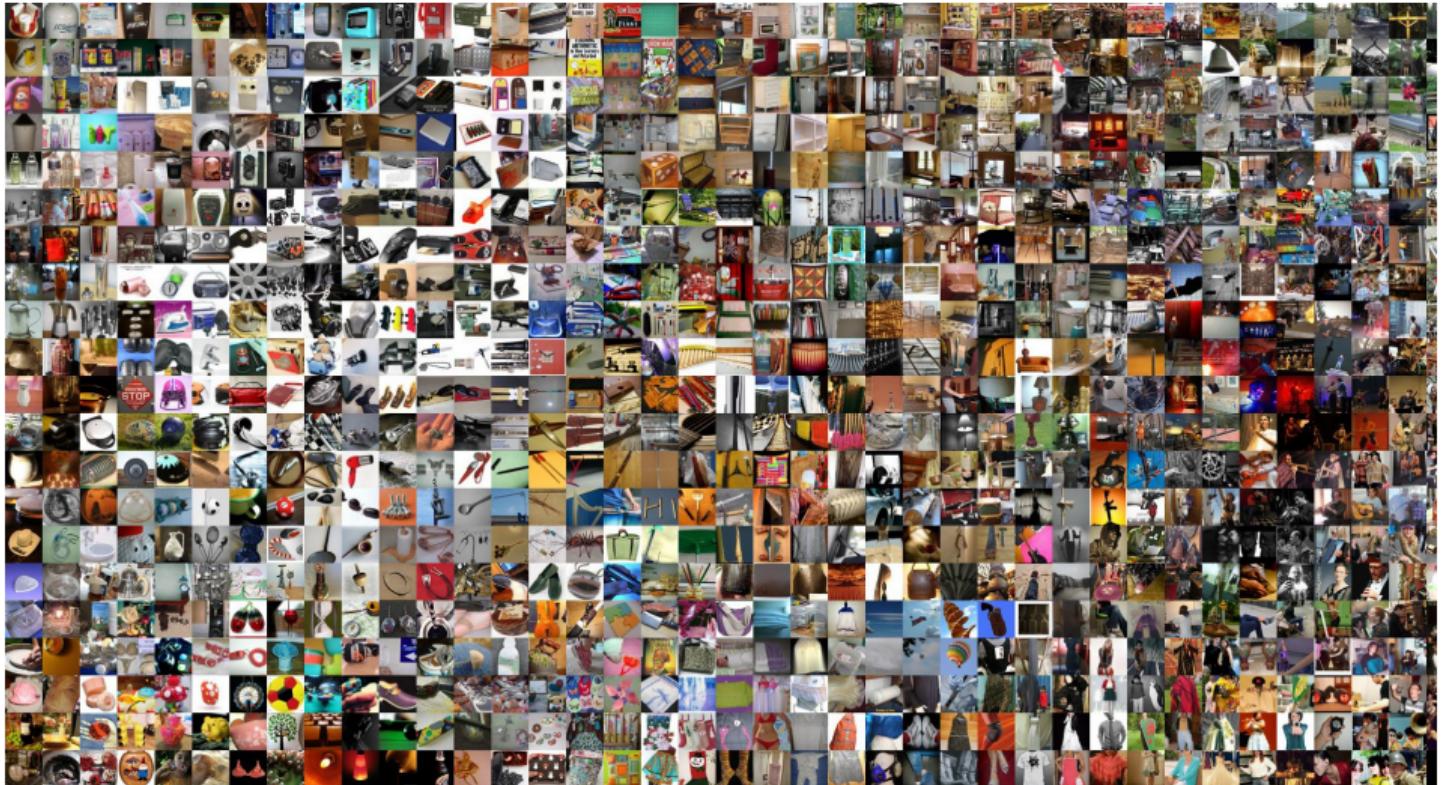
# Computer Graphics vs. Computer Vision



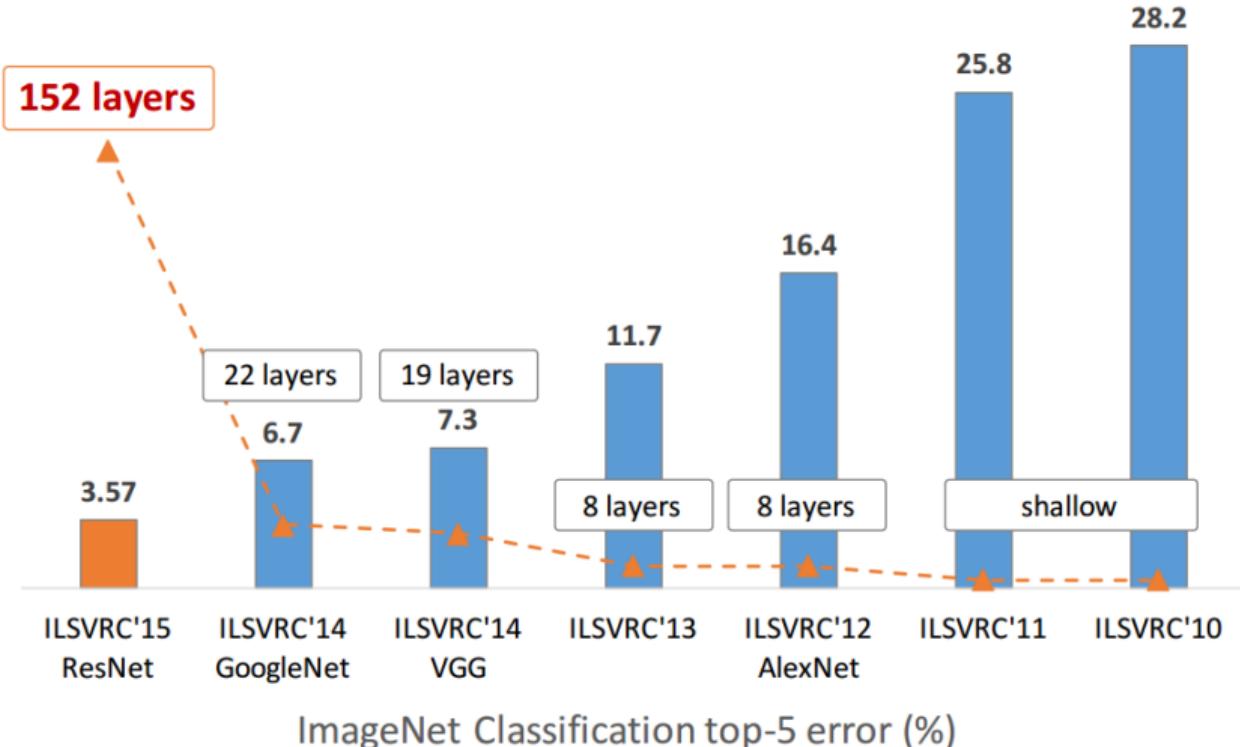
Computer Graphics can also be an ill-posed problem:

- ▶ Rendering is slow, fill-in missing information
- ▶ Deblurring, denoising, superresolution

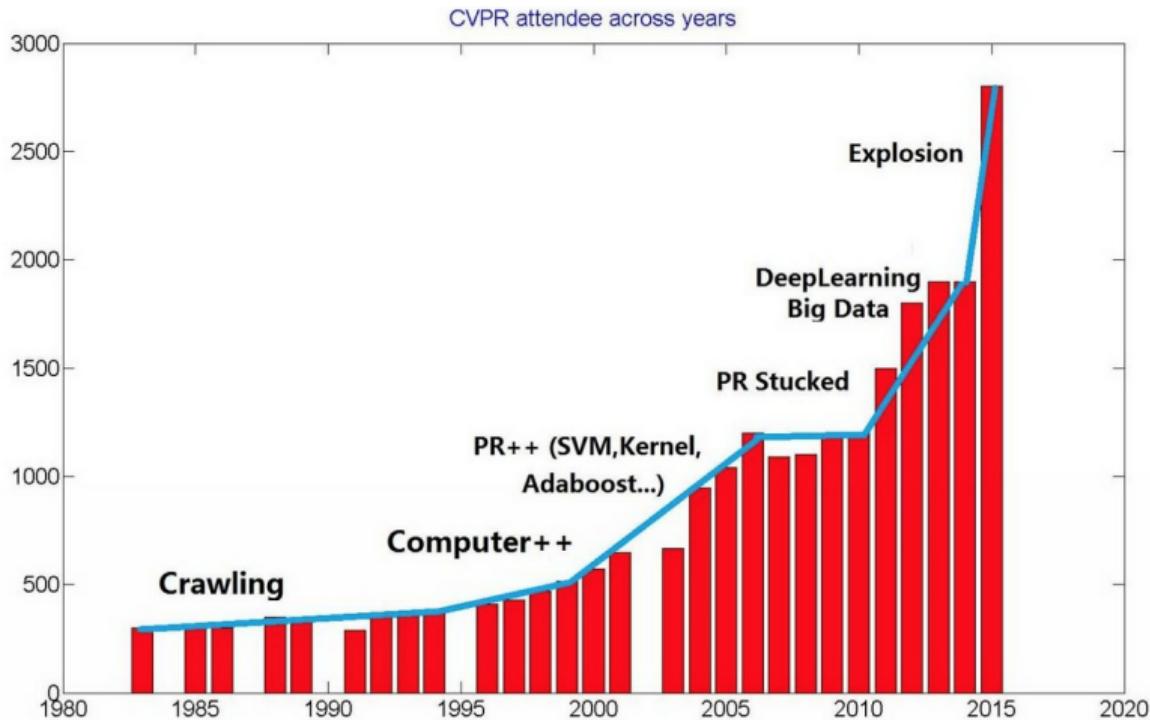
# What is all the fuss about?



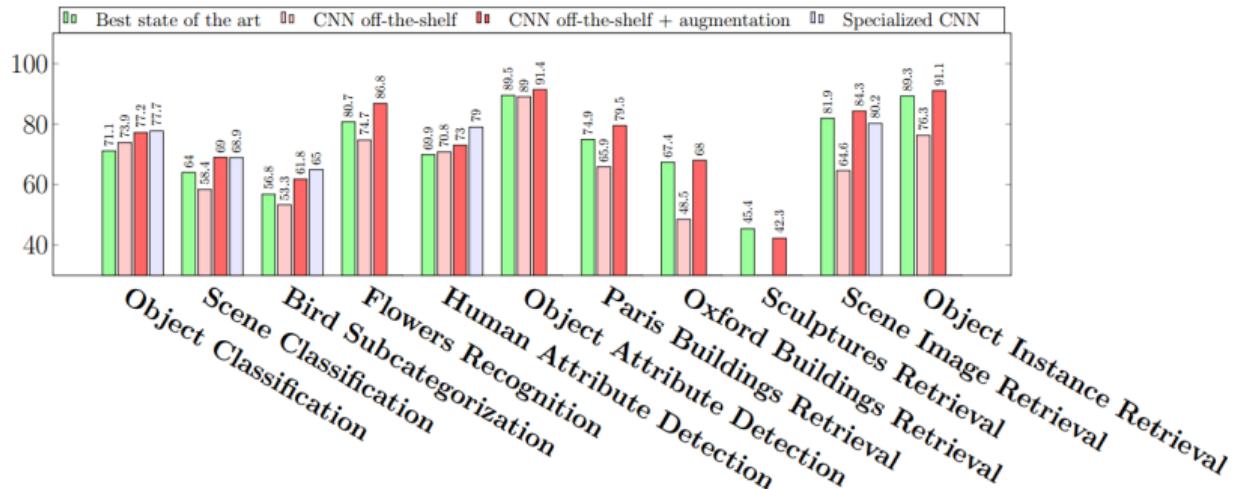
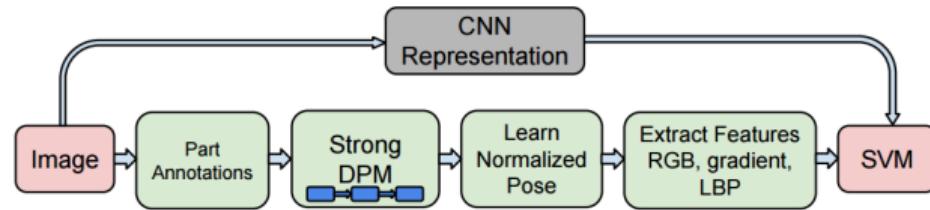
# The Deep Learning Revolution



# CVPR Attendance



# Generalization



[Razavian, Arxiv 2014]

# AI Startups



## 100 STARTUPS USING ARTIFICIAL INTELLIGENCE TO TRANSFORM INDUSTRIES

### CONVERSATIONAL AI/ BOTS



### VISION



### AUTO



### ROBOTICS



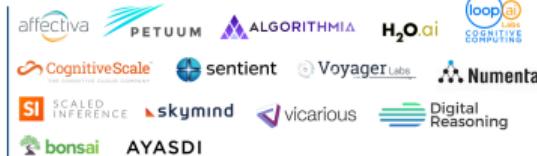
### CYBERSECURITY



### BUSINESS INTELLIGENCE & ANALYTICS



### CORE AI



### AD, SALES, CRM



### HEALTHCARE



### FINTECH & INSURANCE



### OTHER



### TEXT ANALYSIS/ GENERATION



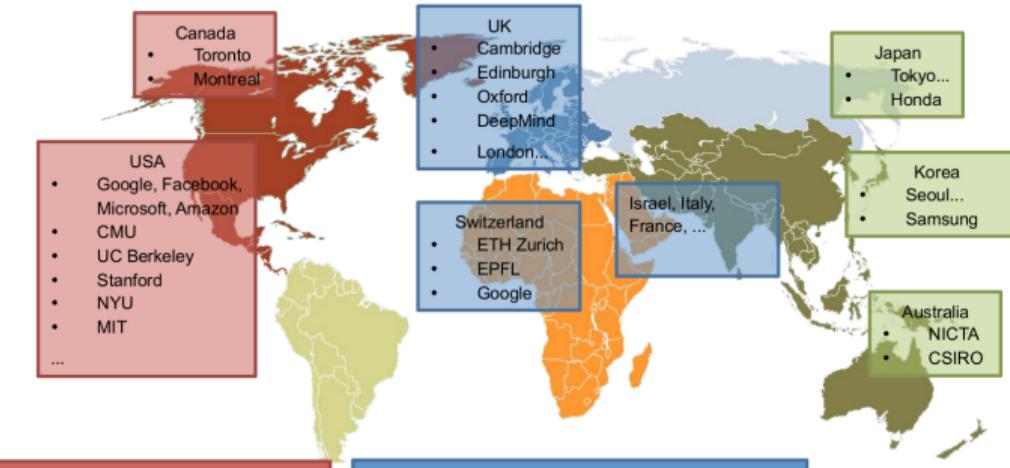
### IOT/IIOT



### COMMERCE



# AI Research Worldwide



## Facebook AI Research

- New lab founded in 2013
- New arm in Paris: 2015
- Europe University partnership program 02/2016

## Google DeepMind (London)

- Founded in 2011 as AI research startup
- 2014 acquisition through Google (\$0.4bn)
- 02/2016: DeepMind learning algorithm beats European Go champion

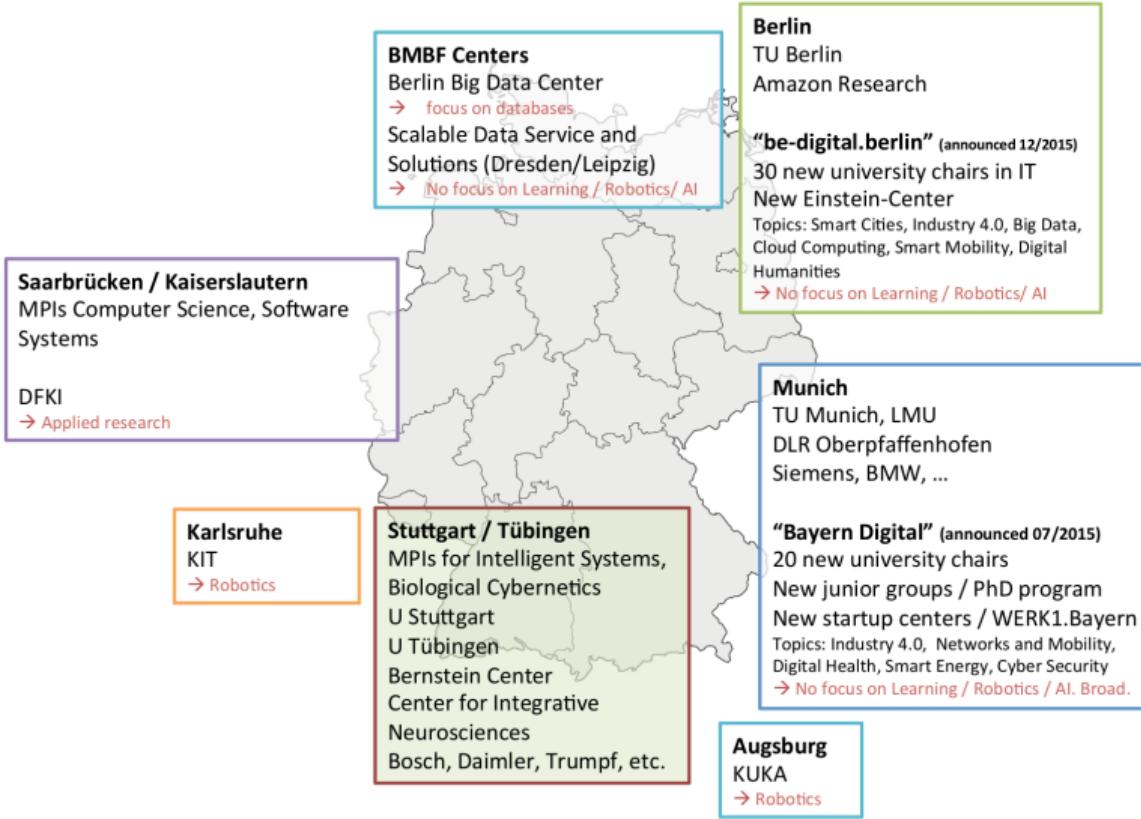
## Toyota Research Institute (Palo Alto)

- 11/2015 Toyota announced that it is investing \$1bn in the next 5 years
- New AI and robotics R&D arm in Silicon Valley
- Two research centers at Stanford and MIT (\$50m in 5 years)

## OpenAI (San Francisco)

- Non profit AI research company founded in 12/2015
- \$1bn pledged from Elon Musk, Amazon and others

# AI Research in Germany



# Cyber Valley

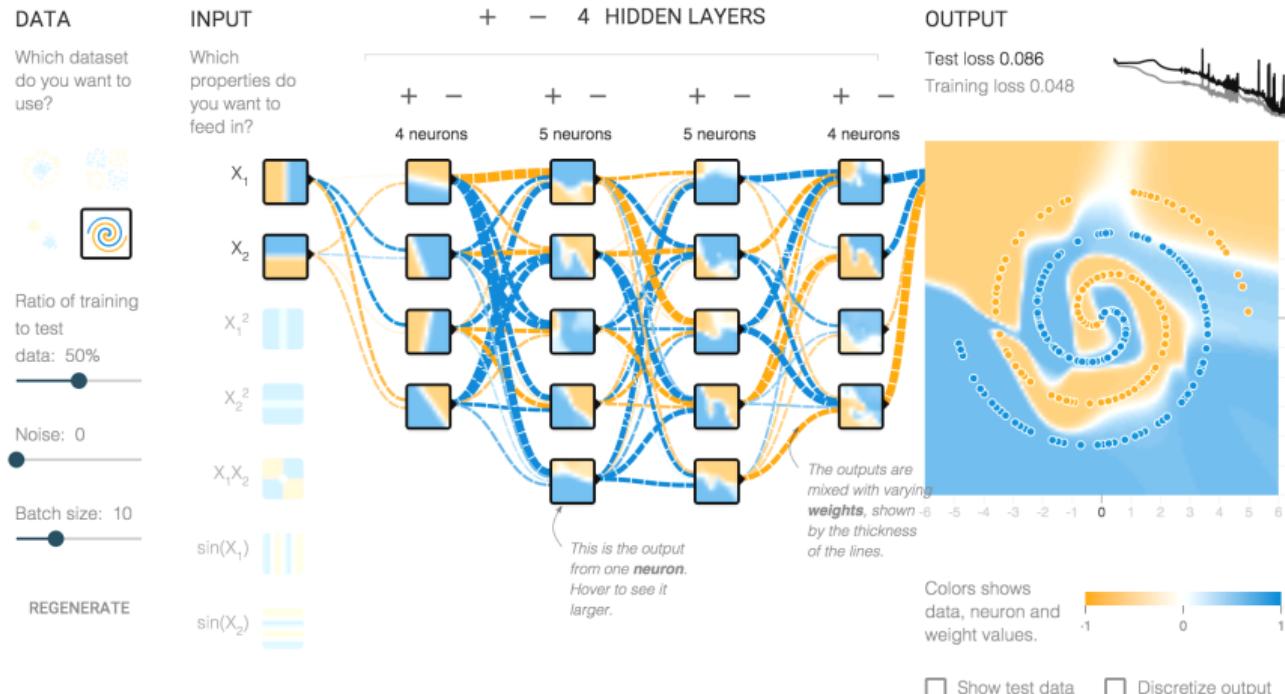


- ▶ Internationally visible hub for research on AI & IS
- ▶ Max Planck, universities, government and industry
- ▶ Foster entrepreneurship, educate engineers for the IS economy, IMPRS

Why is the AI revolution happening now?

# #1 Machine Learning Starts Working

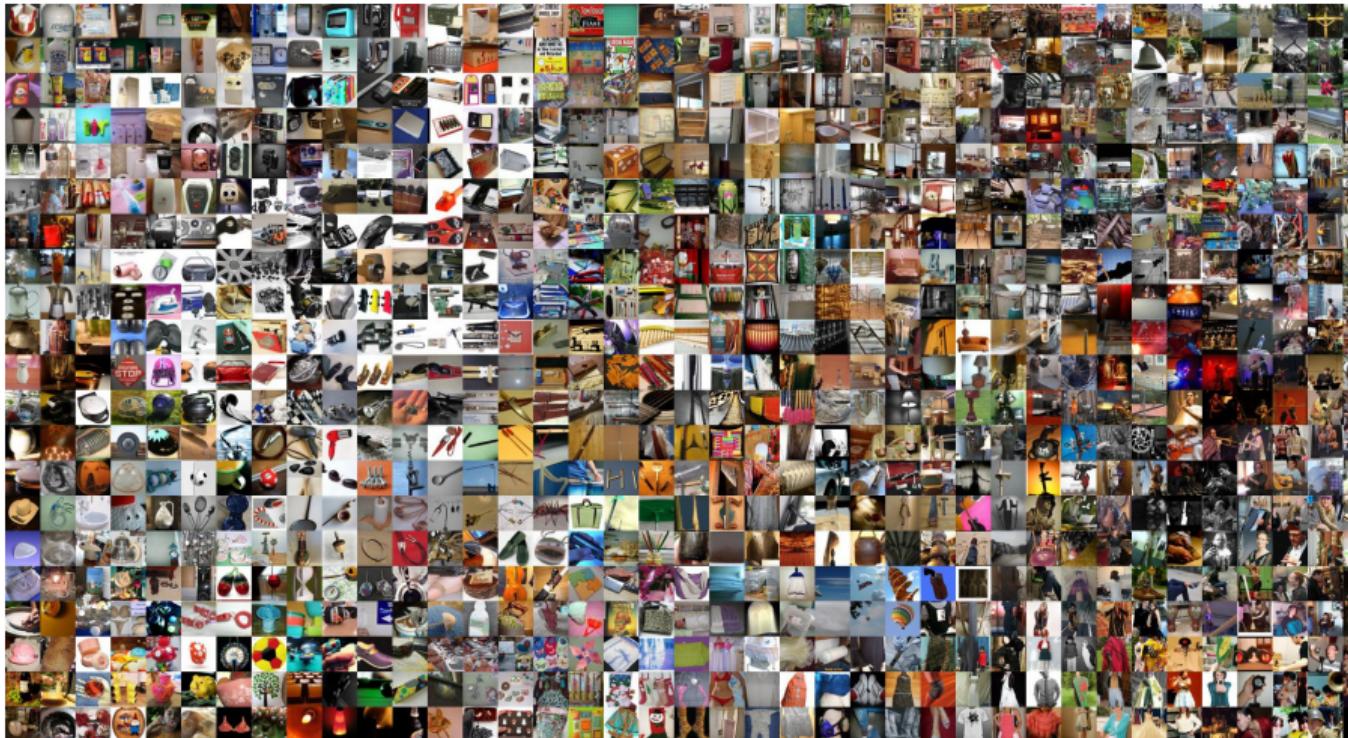
# Machine Learning Starts Working



[playground.tensorflow.org](http://playground.tensorflow.org)

#2 Big Data

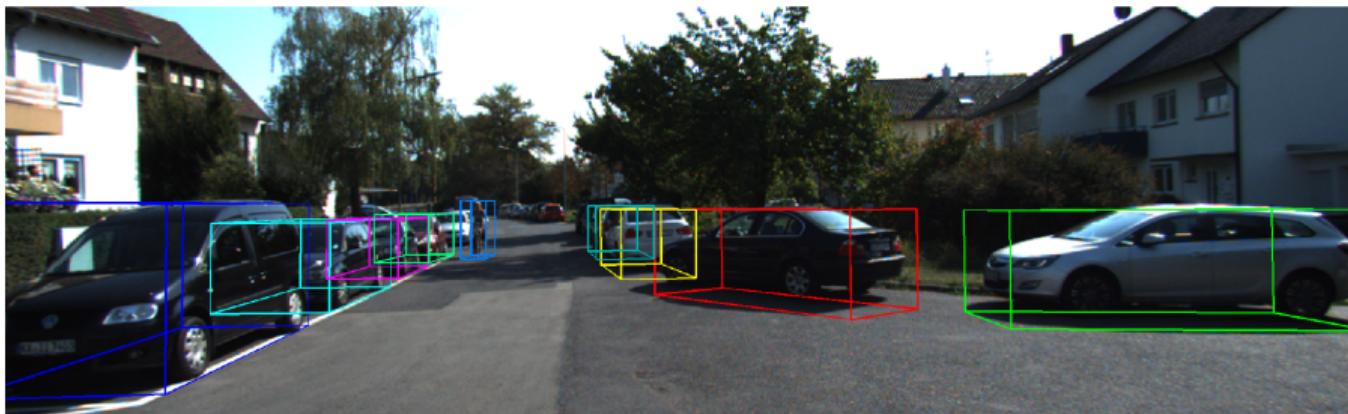
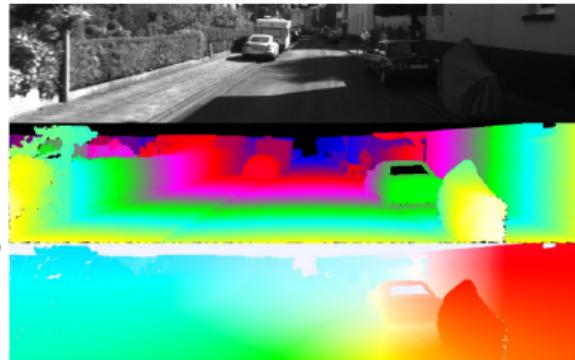
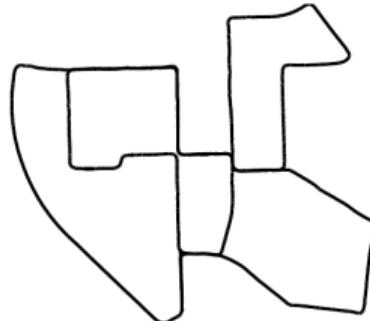
# Big Data



[www.image-net.org](http://www.image-net.org)

# Big Data

360° Velodyne Laserscanner  
Stereo Camera Rig  
GPS



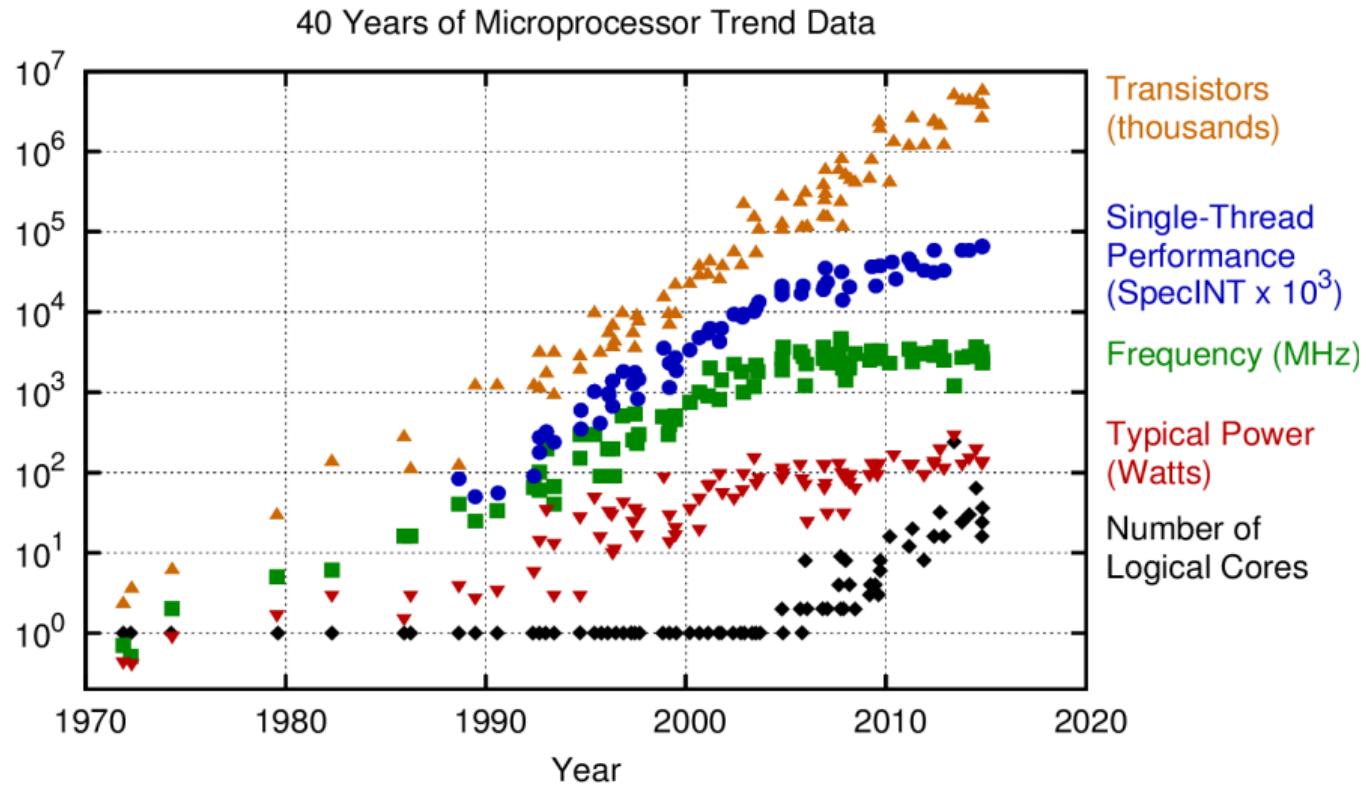
# Big Data



[www.cityscapes-dataset.com](http://www.cityscapes-dataset.com)

#3 Moore's Law

# Moore's Law



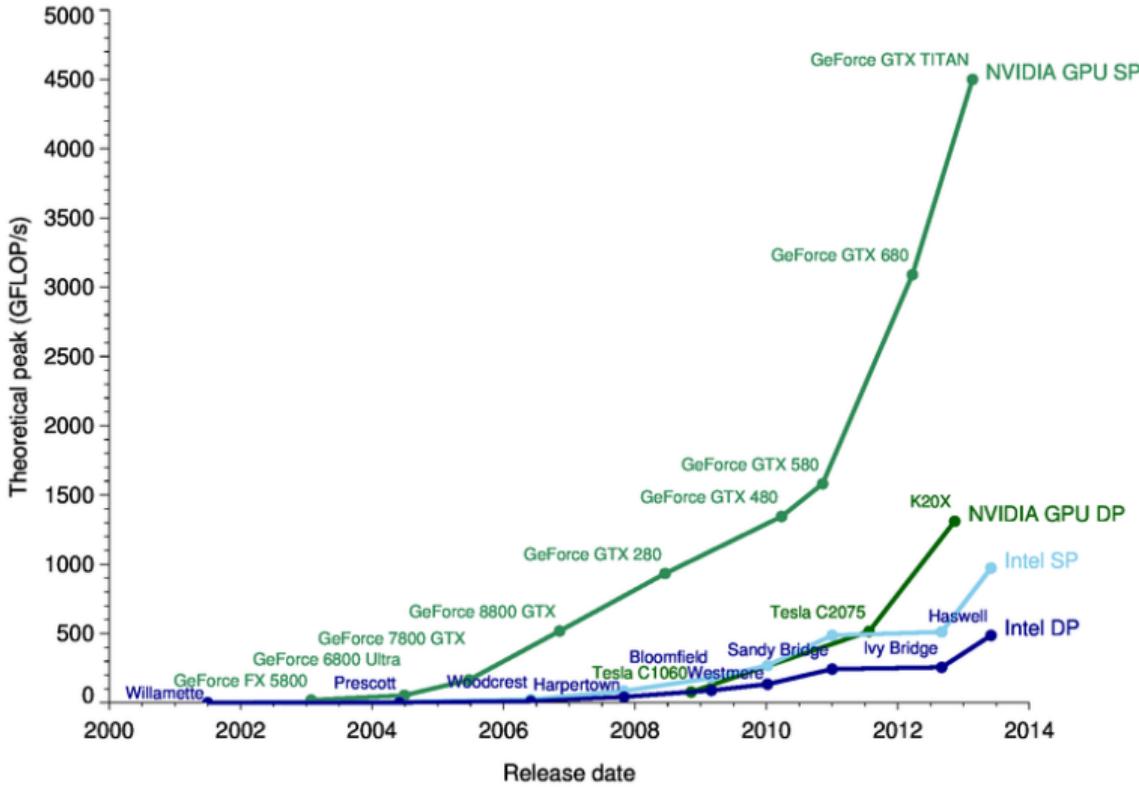
Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2015 by K. Rupp

# Moore's Law



[www.nvidia.com](http://www.nvidia.com)

# Moore's Law



## #4 Coorporate Research Philosophy

# Coorporate Research Philosophy

- ▶ Geoffrey Hinton → Google
- ▶ Yann LeCun → Facebook
- ▶ Andrew Ng → Baidu
- ▶ Zoubin Ghahramani → Uber
- ▶ Ian Goodfellow → OpenAI
- ▶ Marc Pollefeys → Microsoft
- ▶ Trevor Darrell → Nexar
- ▶ Fei-Fei Li → Google
- ▶ Marc'Aurelio Ranzato → Google
- ▶ Alex Smola → Amazon
- ▶ David Nister → Tesla
- ▶ Raquel Urtasun → Uber
- ▶ Pushmeet Kohli → Google

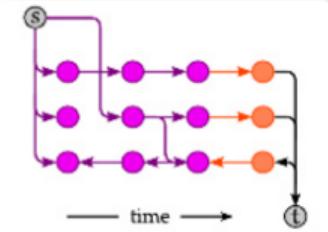
#5 Open Access

# Open Access



### 3D Reconstruction

Probabilistic volumetric 3D reconstruction with ray potentials.



### FollowMe

Computationally Efficient Online Min-Cost Flow Tracking.



### Discrete Flow

Optical flow estimation cast as inference in a discrete CRF.

$$f(x) = \begin{cases} \xi(x) & \text{if } x \in \mathcal{S} \\ 0 & \text{otherwise} \end{cases}$$
$$\xi(x) : D^N \rightarrow \mathbb{R}$$
$$|\mathcal{S}| \ll 2^N$$

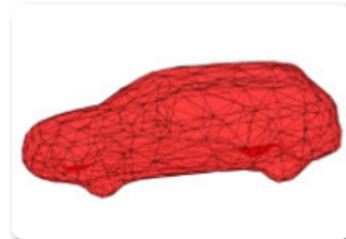
### LIBSMS

Efficient MAP inference with sparse high-order potentials.



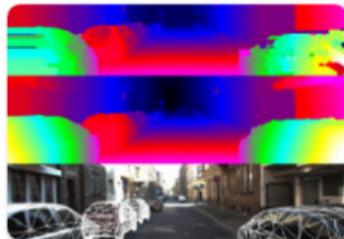
### Indoor Scenes

3D layout and 3D object inference from a single RGB-D image.



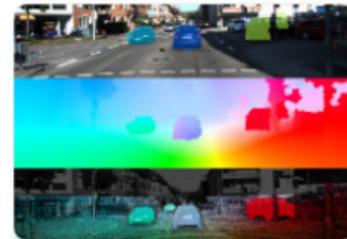
### Semi-Convex Hull

Reduces a complex CAD model to a watertight low-face mesh.



### Displets

Resolving stereo ambiguities using knowledge about objects.



### Object Scene Flow

Scene flow for autonomous vehicles. Novel realistic dataset.

# Open Access



## #6 Rapid Dissemination

# Rapid Dissemination

The screenshot shows a Cornell University Library watermark at the top left. At the top right, there is a message of thanks to the Simons Foundation and The Alliance of Science Organisations in Germany, coordinated by TIB, MPG and HGF. The main navigation bar includes links for 'arXiv.org > cs > arXiv:1701.04722', a search bar ('Search or Article ID inside arXiv'), a dropdown menu ('All papers'), a magnifying glass icon, a link to 'Broaden your search using Semantic Scholar', and another search bar.

The page title is 'Computer Science > Learning'. The main title of the paper is 'Adversarial Variational Bayes: Unifying Variational Autoencoders and Generative Adversarial Networks'. The authors are Lars Mescheder, Sebastian Nowozin, and Andreas Geiger. The submission date is 'Submitted on 17 Jan 2017 (v1)', last revised '8 Mar 2017 (this version, v2)'. The abstract discusses Variational Autoencoders (VAEs) and Generative Adversarial Networks (GANs), stating that their method, Adversarial Variational Bayes (AVB), provides a principled connection between VAEs and GANs, achieving exact maximum-likelihood assignment for generative model parameters and exact posterior distributions over latent variables.

On the right side, there is a 'Download' section with links for PDF and other formats, and a 'License' link. It also shows the current browse context as 'cs.LG' with links for '< prev | next >', 'new | recent | 1701', and a 'Change to browse by' dropdown set to 'cs'. Below that is a 'References & Citations' section with a NASA ADS link. Further down is a 'DBLP - CS Bibliography' section with links for 'listing | bibtex' and author names: Lars M. Mescheder, Sebastian Nowozin, and Andreas Geiger. The bottom right features a 'Bookmark' section with various sharing icons.

[Mescheder et al., Arxiv 2017]

# Rapid Dissemination

The screenshot shows a GitHub Gist page for a notebook named "Adversarial variational bayes toy example.ipynb". The page includes a header with "GitHubGist" and a search bar, navigation links for "All gists" and "GitHub", and buttons for "New gist" and "Edit". Below the header, the gist title is displayed along with the author's profile picture and the last active time ("Last active 4 days ago"). The notebook has 39 stars and 12 forks. There are buttons for "Code", "Revisions 2", "Stars 39", "Forks 12", "Embed", "Raw", and "Download ZIP". The notebook content starts with a section titled "Adversarial Variational Bayes toy example" by Ben Poole (January 26, 2017). It describes the implementation of a toy example from the paper "Adversarial Variational Bayes: Unifying Variational Autoencoders and Generative Adversarial Networks" by Lars Mescheder, Sebastian Nowozin, and Andreas Geiger. It also mentions a blog post for another derivation and implementation.

**Adversarial variational bayes toy example.ipynb**

Last active 4 days ago

Code Revisions 2 Stars 39 Forks 12 Embed Raw Download ZIP

**Adversarial Variational Bayes toy example.ipynb**

**Adversarial Variational Bayes toy example**

Ben Poole  
January 26, 2017

This notebook implements the toy example from:  
[Adversarial Variational Bayes: Unifying Variational Autoencoders and Generative Adversarial Networks](#)  
Lars Mescheder, Sebastian Nowozin, Andreas Geiger

See this [blog post](#) for another derivation and implementation of this approach.

**Variational inference**

Given samples  $x$  from an unobserved data density  $q(x)$ , we are interested in building a latent-variable generative model  $p(x, z)$  such that  $p(x)$  is close to  $q(x)$ . This is a difficult problem as evaluating the density under the model,  $p(x)$  requires evaluating an intractable integral over the unobserved latent variable  $z$ :  $p(x) = \int dz p(x, z)$ .

[gist.github.com/poolio/b71eb943d6537d01f46e7b20e9225149](https://gist.github.com/poolio/b71eb943d6537d01f46e7b20e9225149)

# Machine Learning for Graphics and Vision

# Team



Prof. Dr. Hendrik Lensch  
Computer Graphics Group



Prof. Dr. Andreas Geiger  
Autonomous Vision Group

- ▶ TAs:
  - ▶ CGG: Patrick Wieschollek, Arijit Mallick
  - ▶ AVG: Joel Janai, Despoina Paschalidou, Lars Mescheder, Aseem Behl

# Contents

**Goal:** Obtain understanding for basic machine learning concepts and their application to computer vision and computer graphics problems.

## Machine Learning Topics:

- ▶ Classification / Regression
- ▶ kNN, SVM, Random Forest
- ▶ Deep Neural Networks (CNN, RNN)
- ▶ Latent Variable Models (PCA, VAE)
- ▶ Generative Models (GANs)
- ▶ Markov Random Fields

## Vision and Graphics Applications:

- ▶ Object Recognition
- ▶ Semantic Segmentation
- ▶ Reconstruction & Optical Flow
- ▶ Deblurring, Denoising, Superresolution
- ▶ Rendering, Light Transport
- ▶ Global Illumination Sampling

# Organization

- ▶ SWS: 2 V + 2 Ü, 6 ECTS
- ▶ Lecture: Thursdays, 8:15-10:00, F122 (Kleiner Hörsaal), starting April 19
- ▶ Exercise: Fridays, 8:15-10:00, F122 (Kleiner Hörsaal), starting April 20
- ▶ Lectures and exercises will be held in English
- ▶ Oral exam
  - ▶ Dates: 31.7.2018 and 18.9.2018 (mark your calendar!)
  - ▶ 0.3 bonus if 60 % of the points in the exercises
- ▶ Course Website:  
[www.wsi.uni-tuebingen.de/lehrstuhle/computergrafik/lehrstuhl/  
teaching/vorlesung-machine-learning-in-graphics-vision.html](http://www.wsi.uni-tuebingen.de/lehrstuhle/computergrafik/lehrstuhl/teaching/vorlesung-machine-learning-in-graphics-vision.html)
- ▶ **Enroll via ILIAS!**

# Exercises

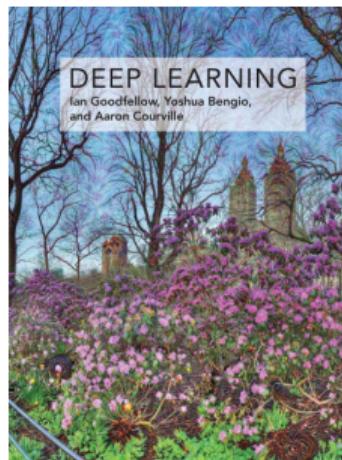
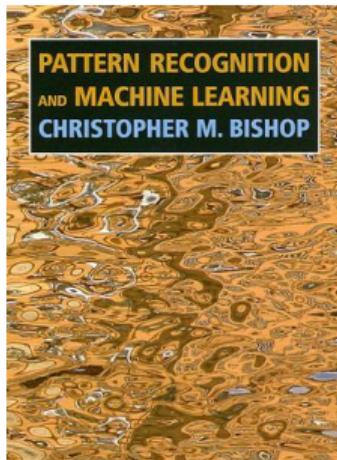
- ▶ Every lecture followed by an exercise
  - ▶ 3 Tutorials: Python, Tensorflow I, Tensorflow II
  - ▶ 6 Assignments (managed via ILIAS, see website), 20 points per exercise
- ▶ Purpose: Introduction to assignments, Q&A, discussion of results
- ▶ Approx. 12 days for completing each assignment
  - ▶ Handout on Friday; Submission until Wednesday 9pm
- ▶ Assignments can be done in groups of up to 2 students
  - ▶ Establish groups in first exercise session (tomorrow!)
  - ▶ Solutions of a group must indicate all group members
  - ▶ Every group member must submit the solution
- ▶ Solutions may not be shared across groups
- ▶ Accessing the computer pool: see course website for instructions

# Lectures

19.04.	01	Machine Learning I
26.04.	02	Machine Learning II
03.05.	03	Deep Learning
17.05.	04	Convolutional Neural Networks
07.06.	05	Advanced Topics in Deep Learning
14.06.	06	Reconstruction & Motion Estimation
21.06.	07	Computer Graphics
28.06.	08	Latent Variable Models
05.07.	09	Generative Models
12.07.	10	Structured Prediction I
19.07.	11	Structured Prediction II
26.07.	12	Deep Structured Models

# Materials

- ▶ Lecture slides, class-room writing (blackboard)
- ▶ Exercise slides & assignments
- ▶ Books:
  - ▶ Bishop: Pattern Recognition and Machine Learning (Springer)
  - ▶ Goodfellow, Bengio, Courville: Deep Learning ([www.deeplearningbook.org/](http://www.deeplearningbook.org/))



Questions?

# Introduction

# Machine Learning



"Field of study that gives computers ability to learn without explicit programming."

Arthur Samuel, 1959

# Machine Learning

"Predicting the future based on past observations."

# Machine Learning



Für größere Ansicht Maus über das Bild ziehen

Asus F540LA-XX274T 39,6 cm (15,6 Zoll) Notebook (Intel core i3-5005U, 8GB Arbeitsspeicher, 1TB Festplatte, Intel HD 5500, Win 10 Home) schwarz  
von [Asus Computer](#)

16 Kundenrezensionen | 38 beantwortete Fragen

Erhältlich bei diesen Anbietern.

Kunden, die diesen Artikel gekauft haben, kauften auch



Microsoft Office 365 Home  
5PCs/MACs - 1  
Jahresabonnement...  
Microsoft

47

Windows 7 / 8 / 10

EUR 73,99



AVM FRITZ!Box 7490  
WLAN AC + N Router  
(VDSL/ADSL, 1.300 Mbit/s (2.4...  
(5 GHz), 450 Mbit/s (2.4...  
 2.460  
EUR 193,90



Tonor 2.4GHz kabellos  
schnurlos Funkmaus  
Wireless Optische Maus  
mit USB 2.0 Empfänger...  
 24  
EUR 8,99



Logitech MK710 Wireless  
Desktop Combo Tastatur  
und Maus (QWERTZ,  
deutsches Tastaturlayout)  
 303  
EUR 73,99

# Machine Learning



## Gummi Saugnapf Glas Kreisschneider

von [Sourcingmap](#)

[Geben Sie die erste Bewertung für diesen Artikel ab](#)

Preis: **EUR 8,98** GRATIS Lieferung innerhalb Deutschlands.

Alle Preisangaben inkl. USt

[Nur noch 12 auf Lager](#)

# Machine Learning



## Gummi Saugnapf Glas Kreisschneider

von [Sourcingmap](#)

Geben Sie die erste Bewertung für diesen Artikel ab

Preis: **EUR 8,98** GRATIS Lieferung innerhalb Deutschlands.

Alle Preisangaben inkl. USt

Nur noch 12 auf Lager

Kunden, die diesen Artikel angesehen haben, haben auch angesehen



Yosoo Flasche  
Fräser/Cutter Glas Flasche  
Cutter Maschine Glas  
Flasche...

★★★★★ 11

EUR 23,99 ✓*Prime*

Sturmhaube Balaclava  
Skimaske 3 Loch Maske  
★★★★★ 69

EUR 2,99 - EUR 7,95

Silverline 282636  
Öl-Glasschneider 175 mm  
★★★★★ 18

EUR 6,88 ✓*Prime*

# Machine Learning

"Show my photos from Utah last August"



"What movies are playing today?"



"Get my call history"

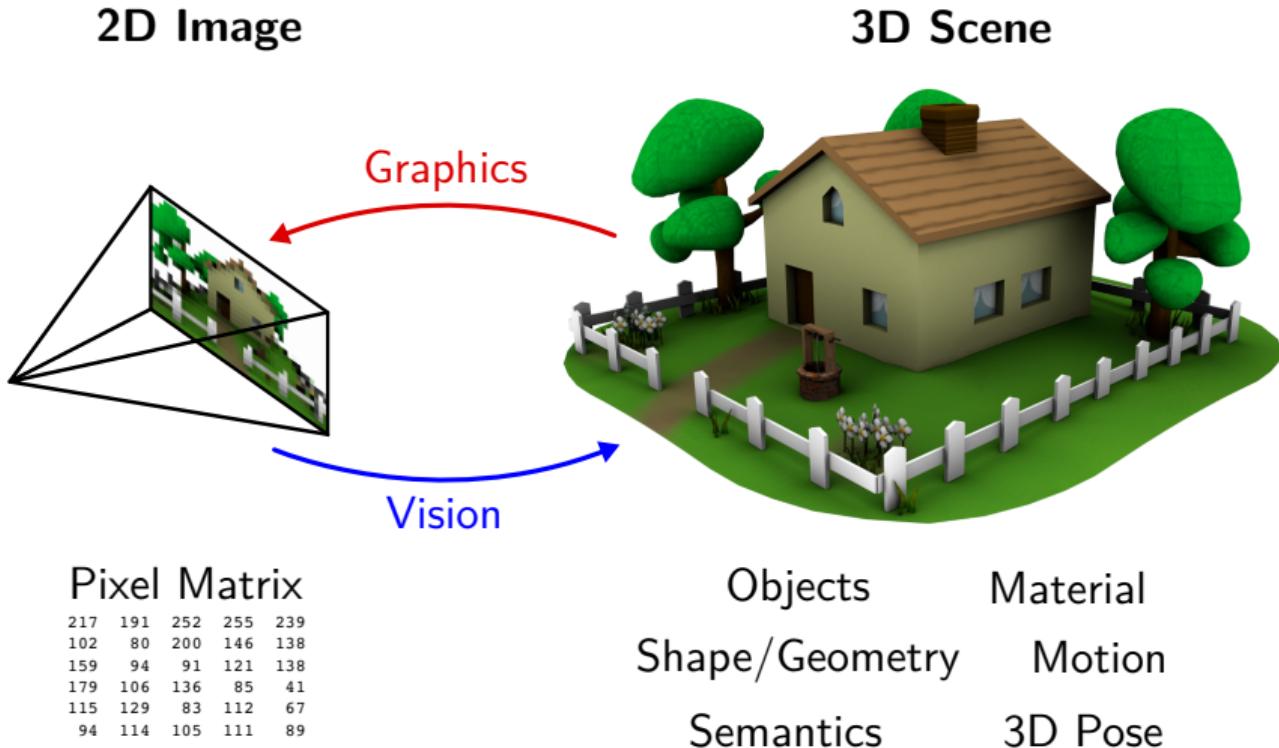


# Machine Learning

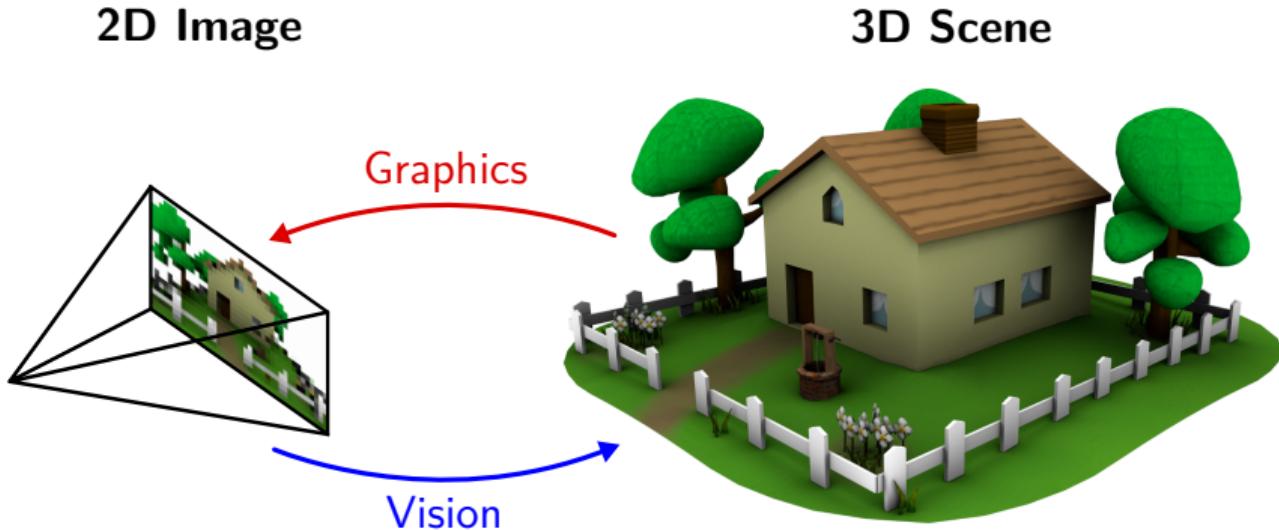


# Computer Vision

# Computer Vision



# Computer Vision



Computer Vision is an ill-posed inverse problem:

- ▶ Many 3D scenes yield the same 2D image
- ▶ Additional constraints (knowledge about world) required

# Why is Computer Vision hard?

MASSACHUSETTS INSTITUTE OF TECHNOLOGY  
PROJECT MAC

Artificial Intelligence Group  
Vision Memo. No. 100.

July 7, 1966

## THE SUMMER VISION PROJECT

Seymour Papert

The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

# Why is Visual Perception hard?



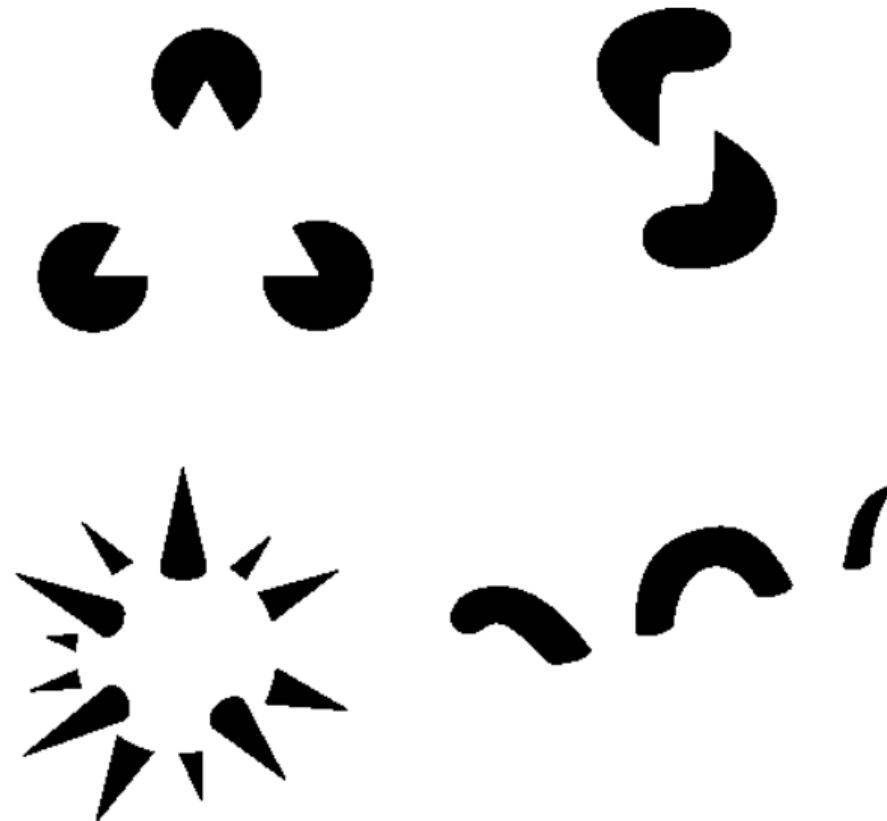
# Why is Visual Perception hard?



# Why is Visual Perception hard?



# Why is Visual Perception hard?

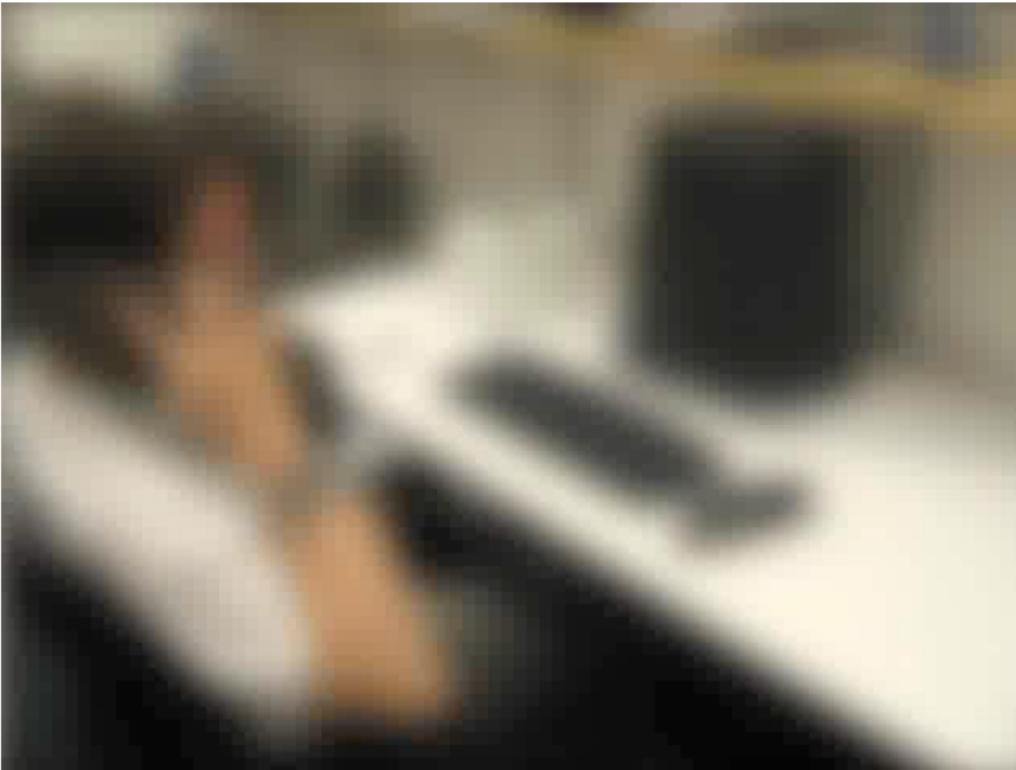


# Why is Visual Perception hard?



Source: <http://www.homeworkshop.com/>

# Why is Visual Perception hard?

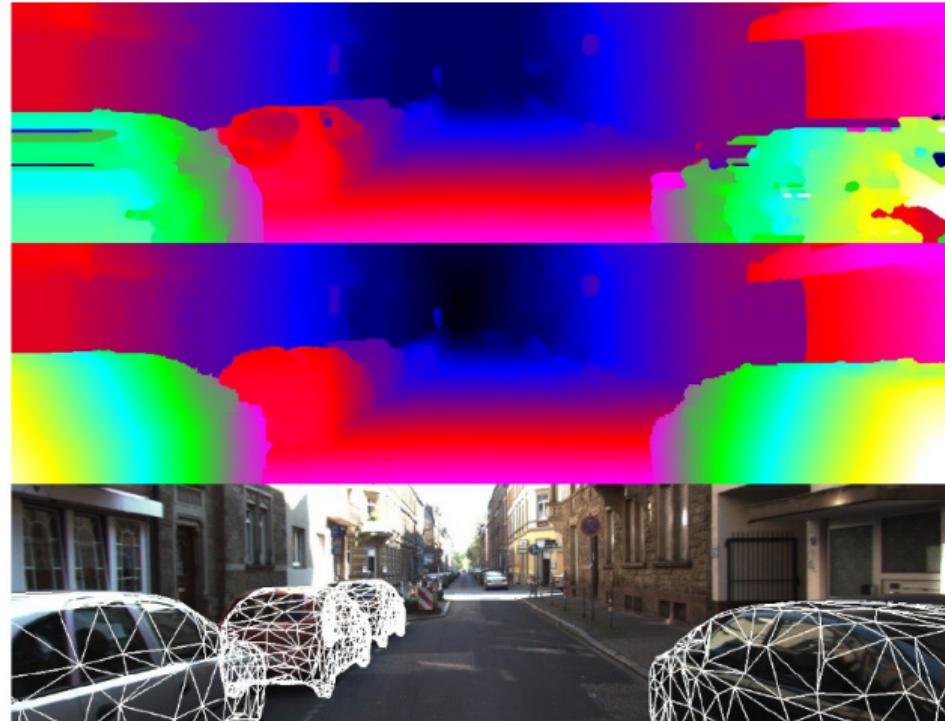


# Why is Visual Perception hard?



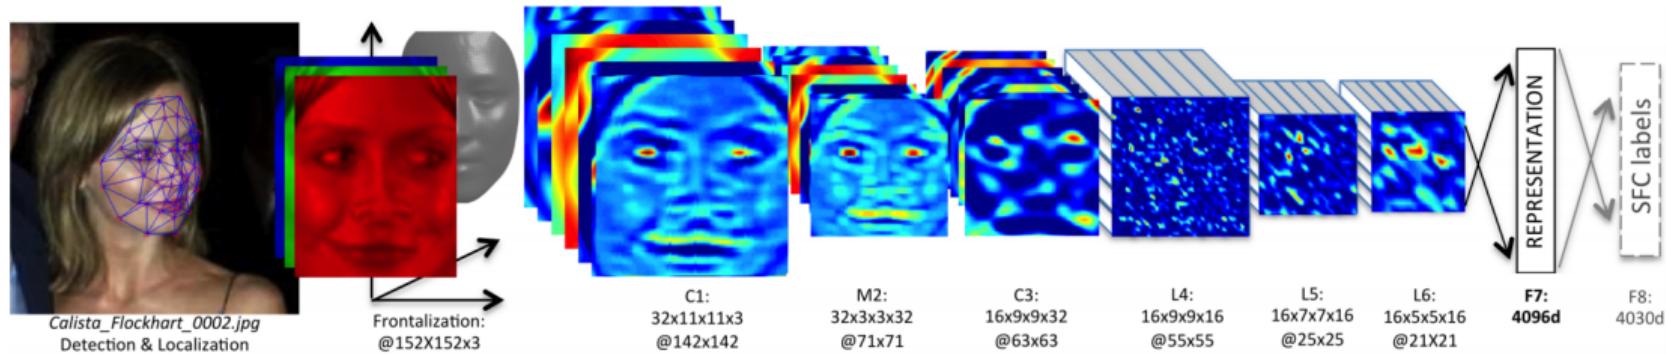
# Some Examples

# Depth Estimation



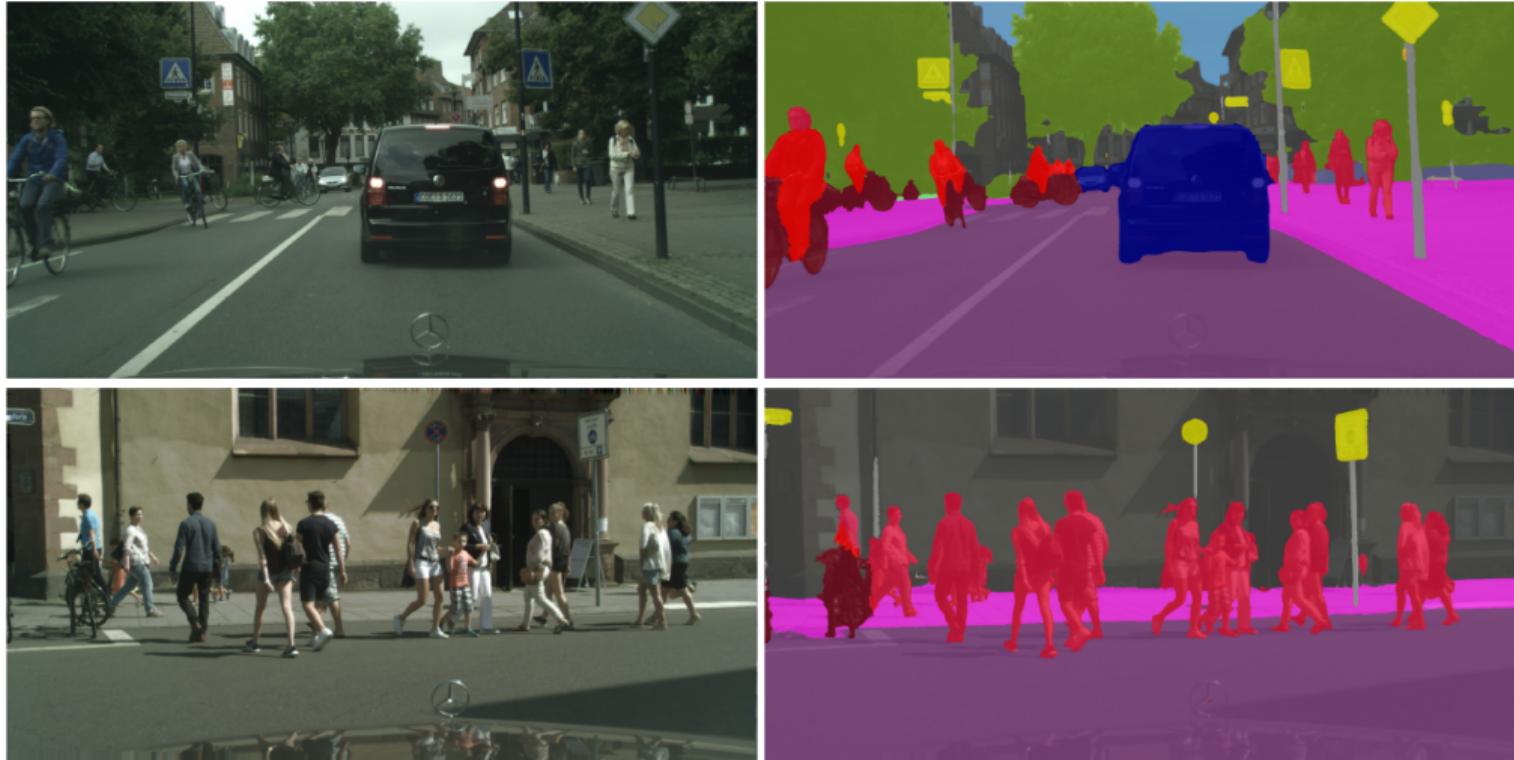
[Güney et al., CVPR 2015]

# Face Recognition



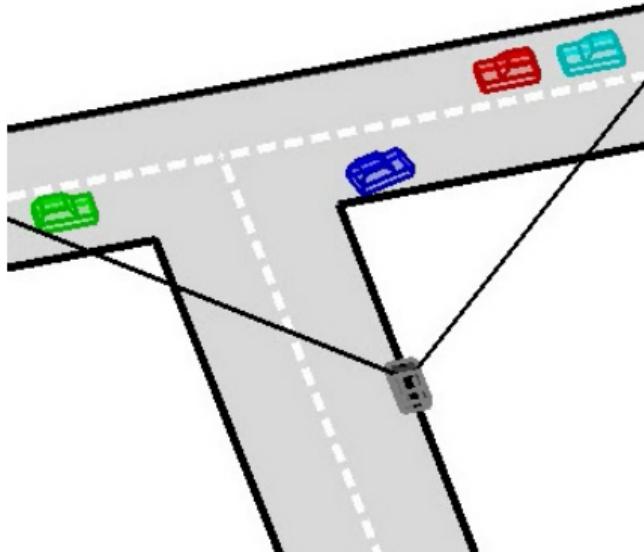
[Taigman et al., CVPR 2014]

# Semantic Segmentation

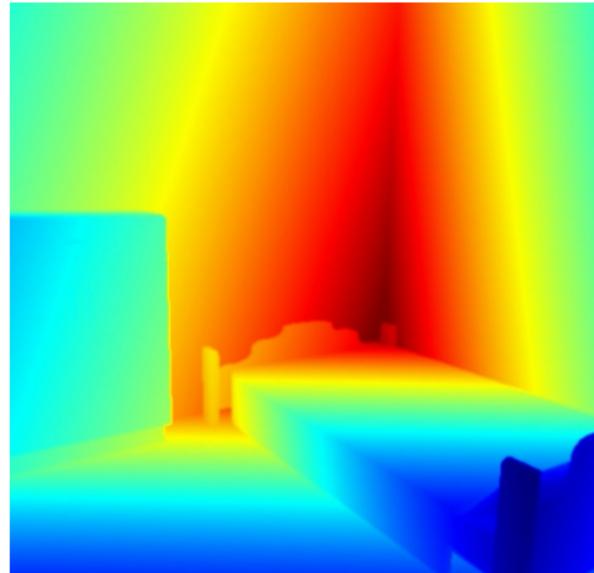


[Kundu et al., CVPR 2016]

# Scene Understanding



[Geiger et al., PAMI 2014]



[Geiger & Wang, GCPR 2015]

# Image Captioning



"little girl is eating piece of cake."



"baseball player is throwing ball in game."



"woman is holding bunch of bananas."



"black cat is sitting on top of suitcase."



"a young boy is holding a baseball bat."



"a cat is sitting on a couch with a remote control."

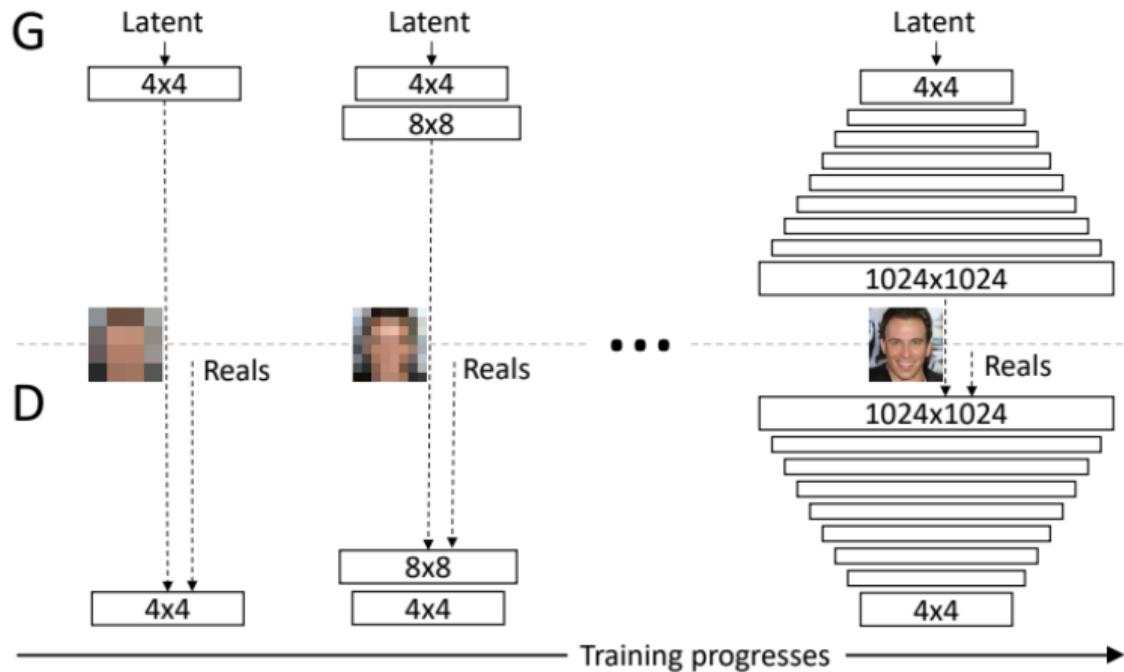


"a woman holding a teddy bear in front of a mirror."



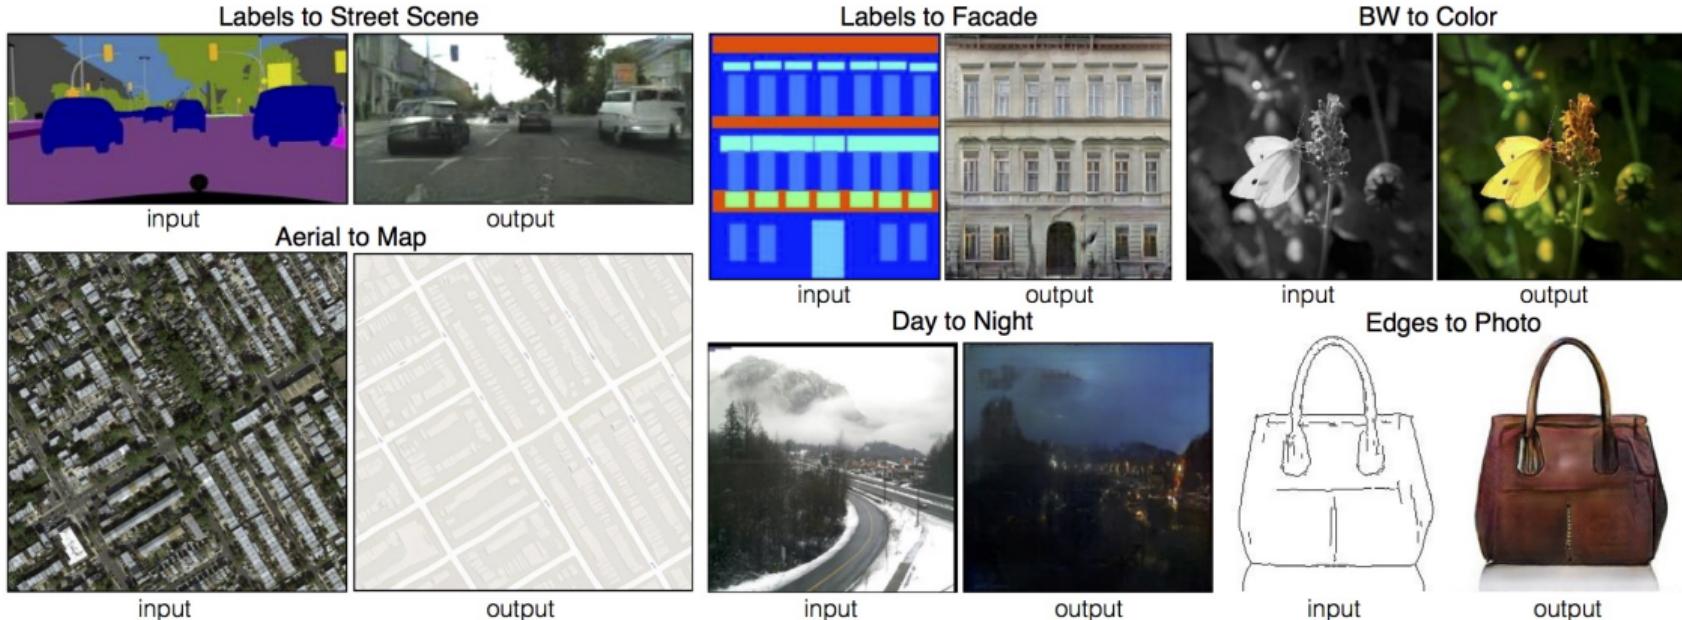
"a horse is standing in the middle of a road."

# Generative Models



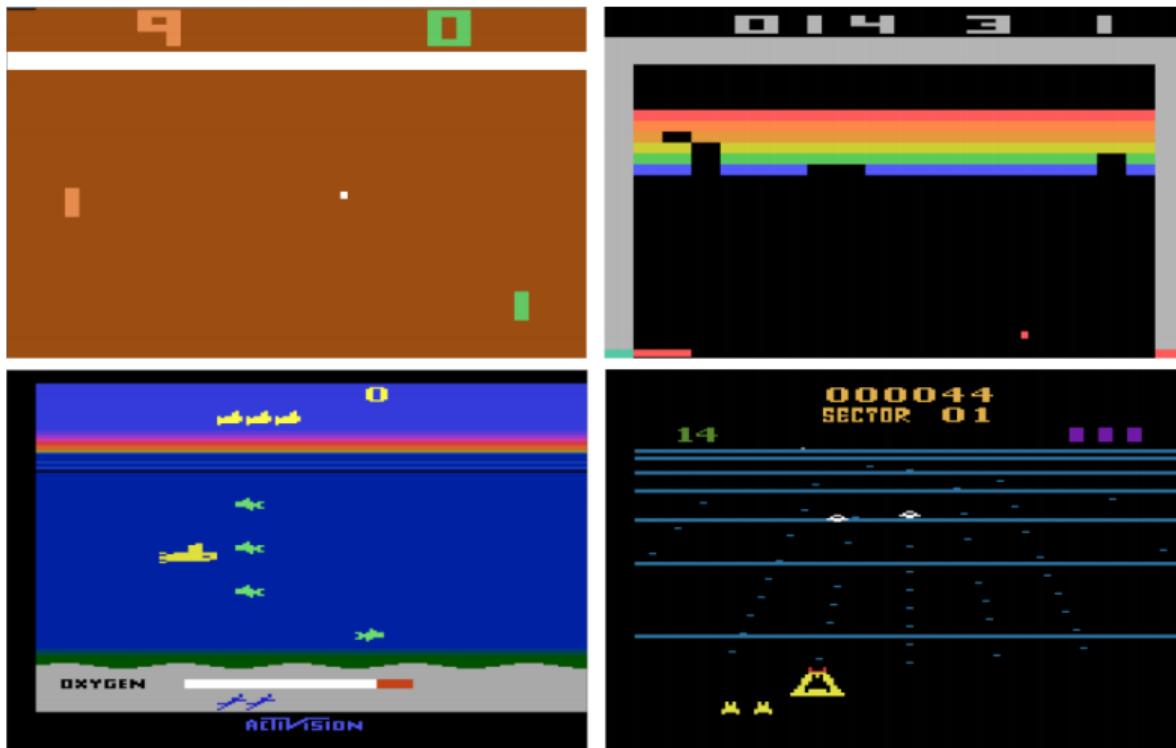
[Karras et al., ICLR 2018]

# Image-to-image Translation



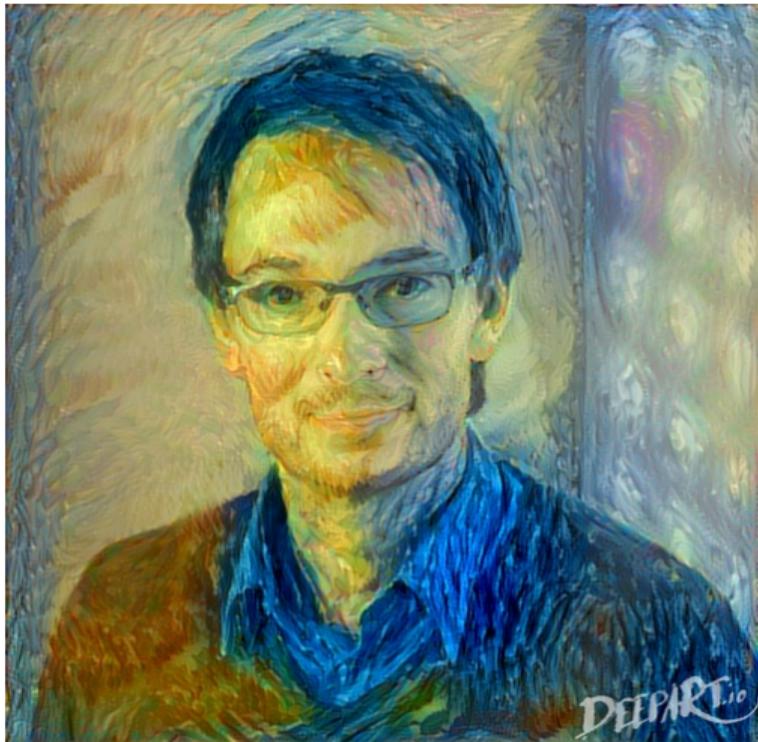
[Isola et al., CVPR 2017]

# Sensori-Motor Control



[Mnih et al., Nature 2015]

# Style Transfer

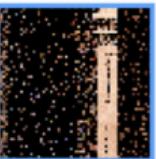


[Gatys et al., CVPR 2016]

# Computer Graphics

# Denoising

(a) 1spp noisy input



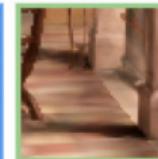
(b) Edge-avoiding wavelets



(c) SURE-based filter



(d) Recurrent autoencoder

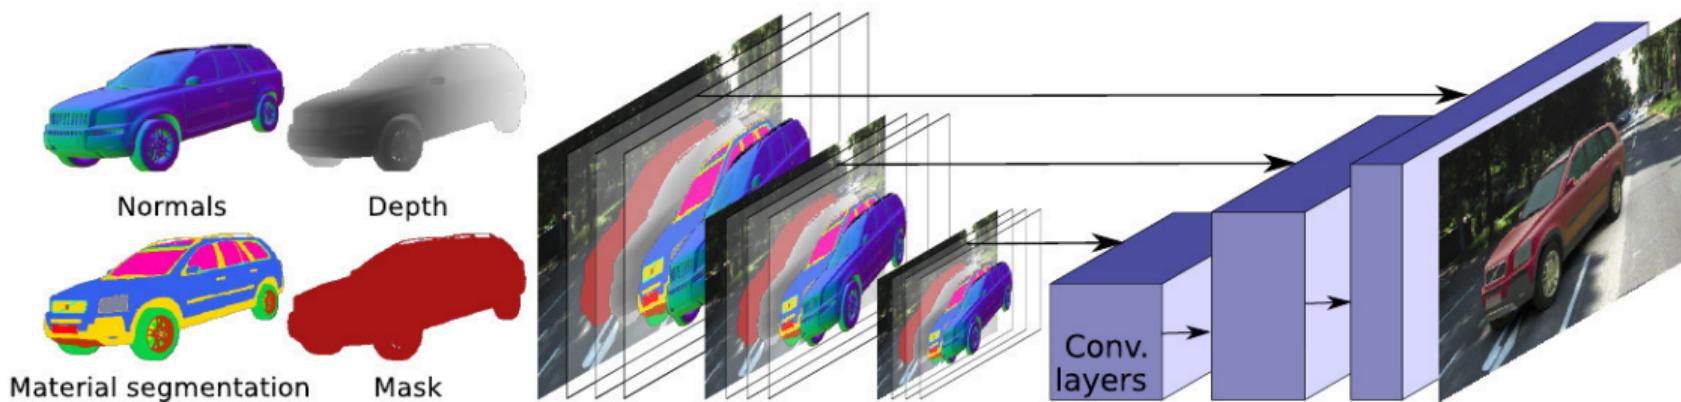


(e) Reference



[Chaitanya et al., SIGGRAPH 2017]

# Rendering



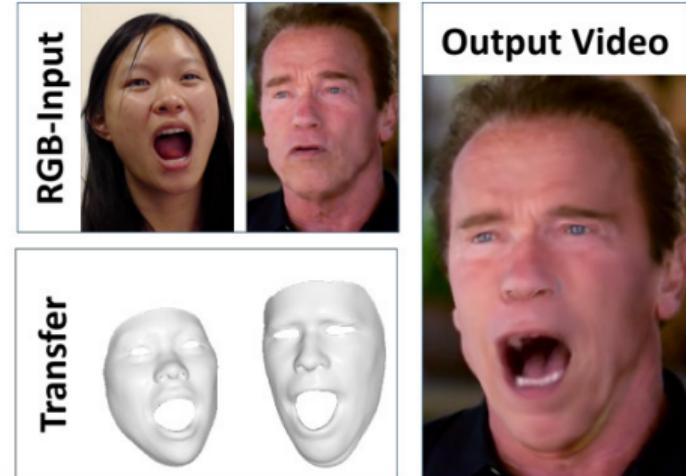
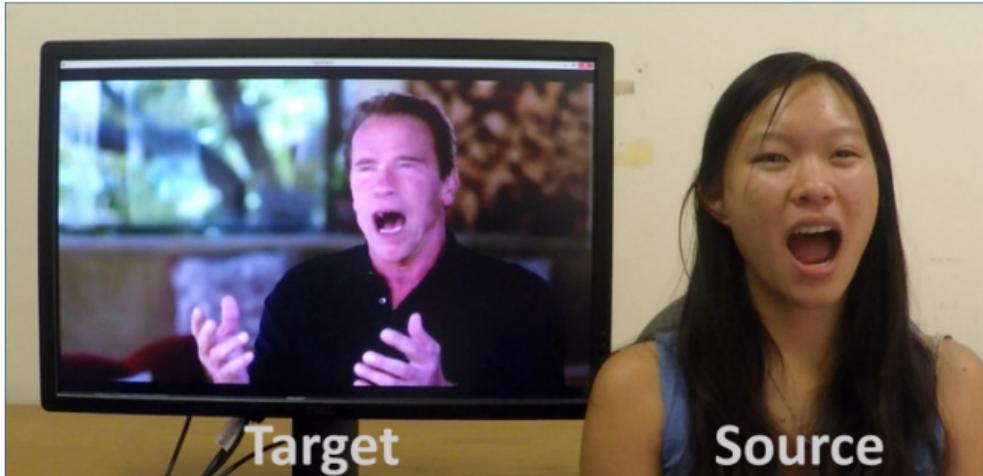
[Alhaija et al., 2018]

# Animation



[Holden *et al.*, SIGGRAPH 2017]

# Reenactment



[Thies et al., CVPR 2016]

# Machine Learning

## Part I

# Machine Learning

## Problems:

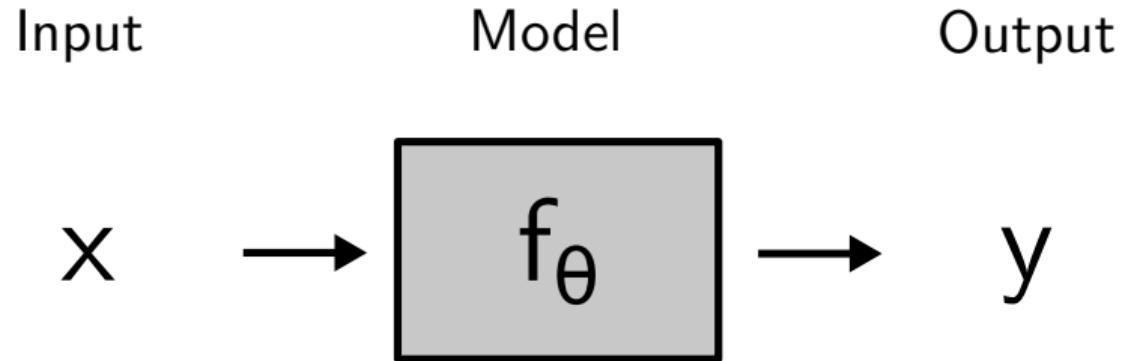
- ▶ Supervised learning
  - ▶ Classification
  - ▶ Regression
  - ▶ Structured prediction
- ▶ Unsupervised learning
  - ▶ Clustering
  - ▶ Density Estimation
  - ▶ Dimensionality reduction
- ▶ Reinforcement learning
- ▶ ...

## Methods:

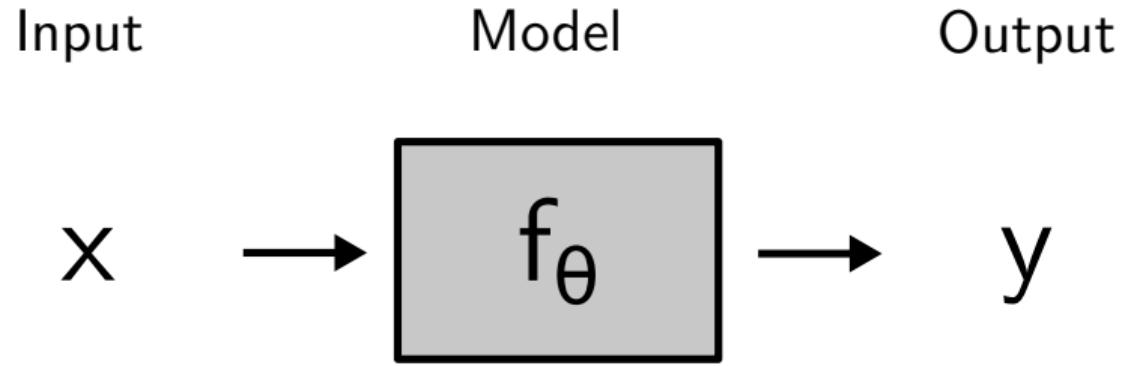
- ▶ Sampling
- ▶ Gaussian Processes
- ▶ Boosting
- ▶ Decision Trees
- ▶ Support Vector Machines
- ▶ Graphical Models
- ▶ Deep Neural Networks
- ▶ ...

# Supervised Learning

# Supervised Learning

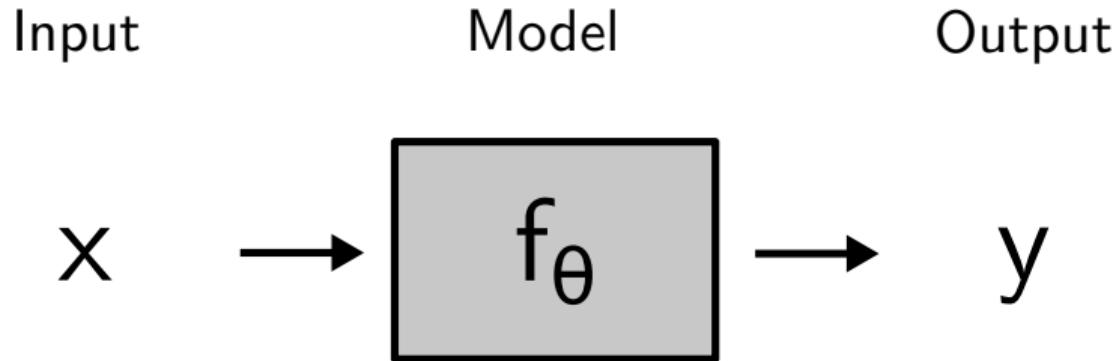


# Supervised Learning



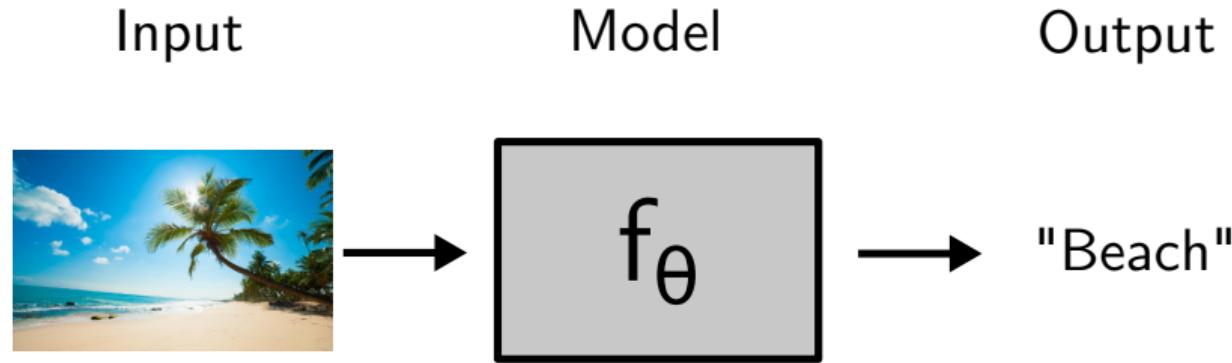
- ▶ **Learning:** Estimate parameters  $\theta$  from training data  $\{(x_i, y_i)\}_{i=1}^N$

# Supervised Learning



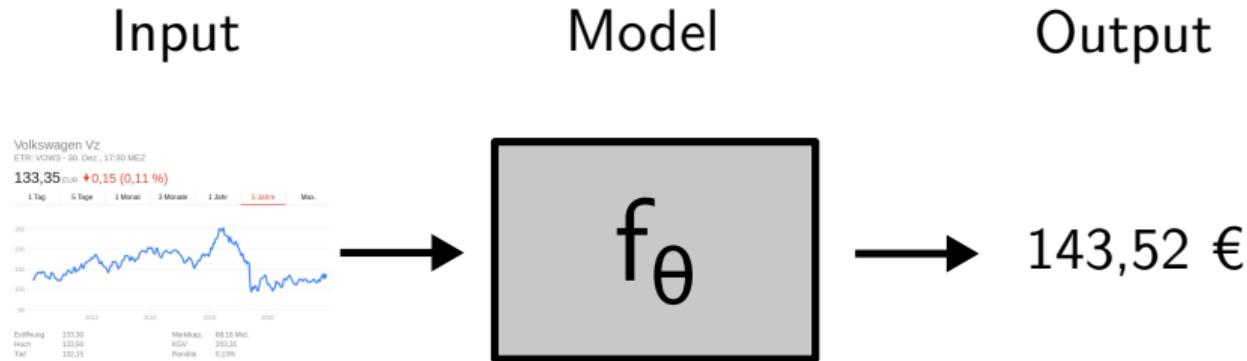
- ▶ **Learning:** Estimate parameters  $\theta$  from training data  $\{(x_i, y_i)\}_{i=1}^N$
- ▶ **Inference:** Make novel predictions:  $y = f_\theta(x)$

# Classification



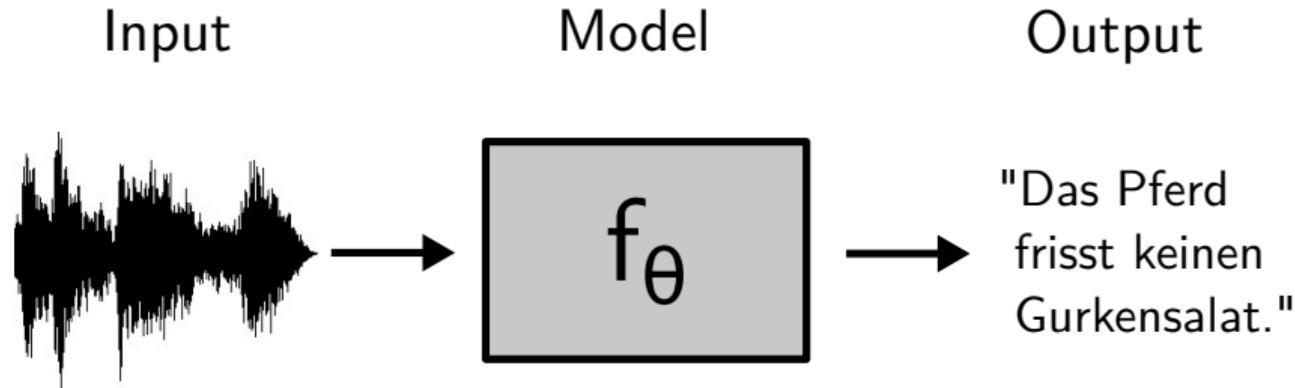
- ▶ **Mapping:**  $f_\theta : \mathbb{R}^{W \times H} \rightarrow \{"\text{Beach}", "\text{No Beach"}\}$

# Regression



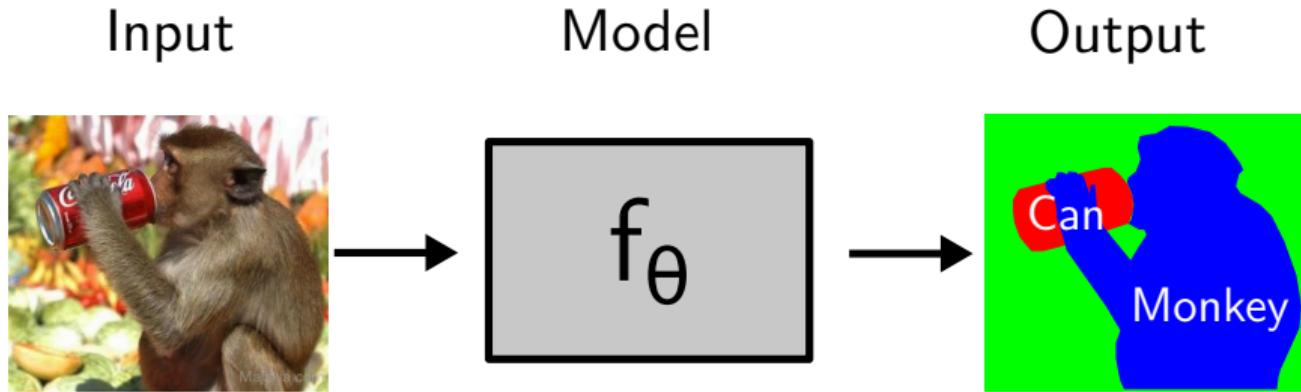
- **Mapping:**  $f_\theta : \mathbb{R}^N \rightarrow \mathbb{R}$

# Structured Prediction



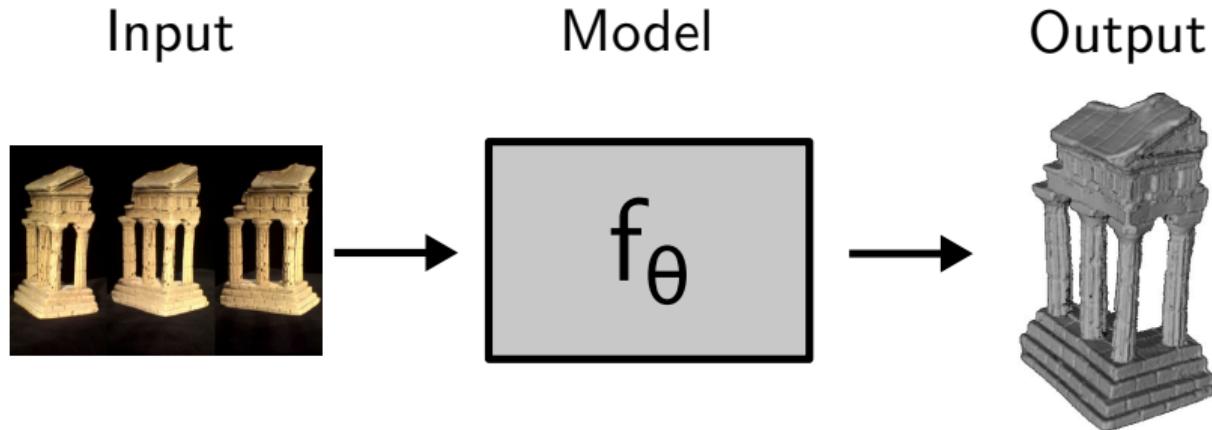
- ▶ **Mapping:**  $f_\theta : \mathbb{R}^N \rightarrow \{1, \dots, L\}^M$

# Structured Prediction



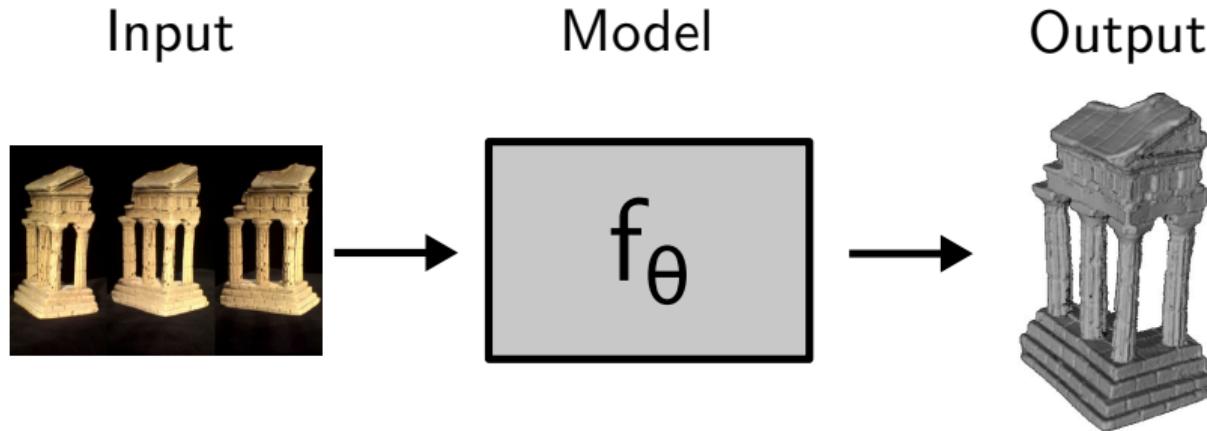
- **Mapping:**  $f_{\theta} : \mathbb{R}^{W \times H} \rightarrow \{1, \dots, L\}^{W \times H}$

# Structured Prediction



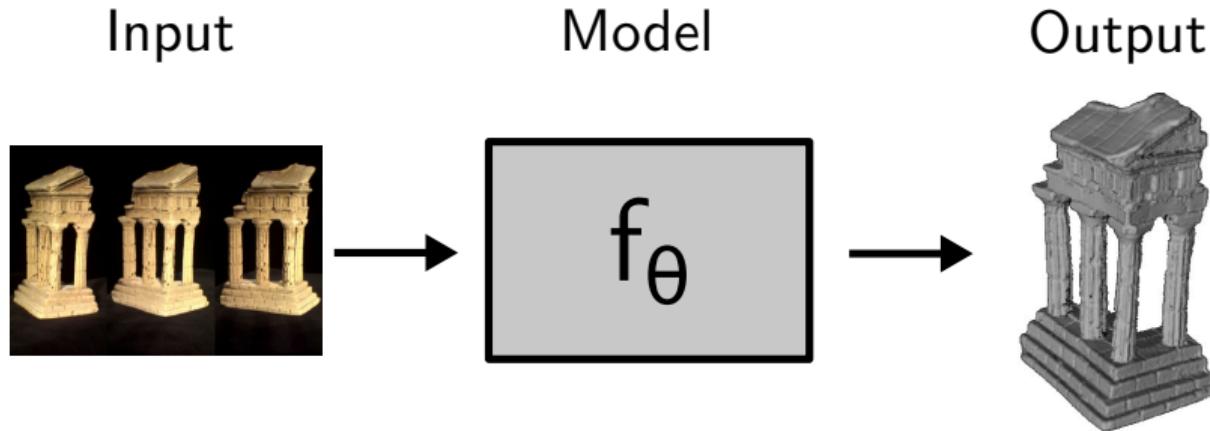
- **Mapping:**  $f_\theta : \mathbb{R}^{W \times H \times N} \rightarrow \{0, 1\}^{M^3}$

# Structured Prediction



- ▶ **Mapping:**  $f_\theta : \mathbb{R}^{W \times H \times N} \rightarrow \{0, 1\}^{M^3}$
- ▶ **Suppose:**  $32^3$  voxels, binary variable per voxel (occupied/free)

# Structured Prediction



- ▶ **Mapping:**  $f_\theta : \mathbb{R}^{W \times H \times N} \rightarrow \{0, 1\}^{M^3}$
- ▶ **Suppose:**  $32^3$  voxels, binary variable per voxel (occupied/free)
- ▶ **Question:** How many different reconstructions?  $2^{32^3} = 2^{32768}$
- ▶ **Comparison:** Number of atoms in the universe?  $\sim 2^{273}$

# Parametric Methods

# Parametric Methods

- ▶ Define parametric model of the observed process / mapping
- ▶ Fit parameters of the model to data
- ▶ Make predictions for new inputs using the parameterized model

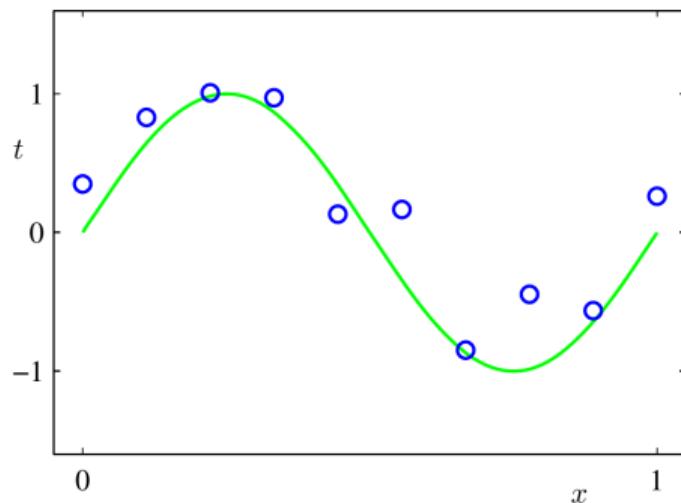
# Polynomial Curve Fitting

# Polynomial Curve Fitting

Let  $\mathcal{X}$  denote a dataset of size  $N$  and let  $(x_n, t_n) \in \mathcal{X}$  denote its elements.

Here  $(x_n, t_n)$  is an input/output pair and we assume  $x_n, t_n \in \mathbb{R}$  for simplicity.

Goal: Predict  $t$  for a previously unseen input  $x$ .



Example with true model (green) and 10 noisy samples drawn from this model (blue).

# Polynomial Curve Fitting

Let us choose a polynomial of order  $M$  to model the dataset  $\mathcal{X}$ :

$$y(x, \mathbf{w}) = \sum_{j=0}^M w_j x^j$$

Tasks:

- ▶ **Training:** Estimate  $\mathbf{w}$  from  $\mathcal{X}$
- ▶ **Inference:** Predict  $t$  for novel  $x$  given estimated  $\mathbf{w}$

# Polynomial Curve Fitting

Let us choose a polynomial of order  $M$  to model the dataset  $\mathcal{X}$ :

$$y(x, \mathbf{w}) = \sum_{j=0}^M w_j x^j$$

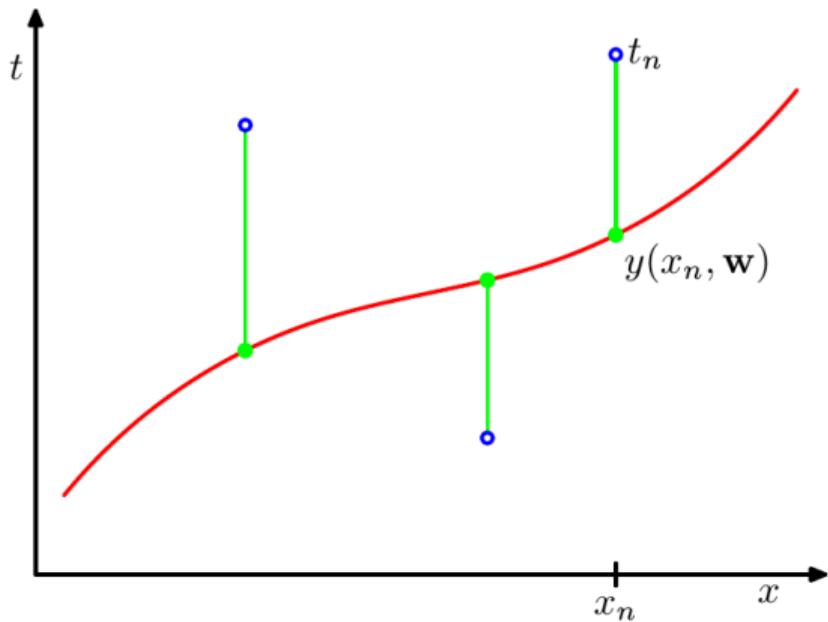
How can we estimate  $\mathbf{w}$  from  $\mathcal{X}$ ?

- ▶ Need to define a loss function / error function to be minimized, e.g.:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2$$

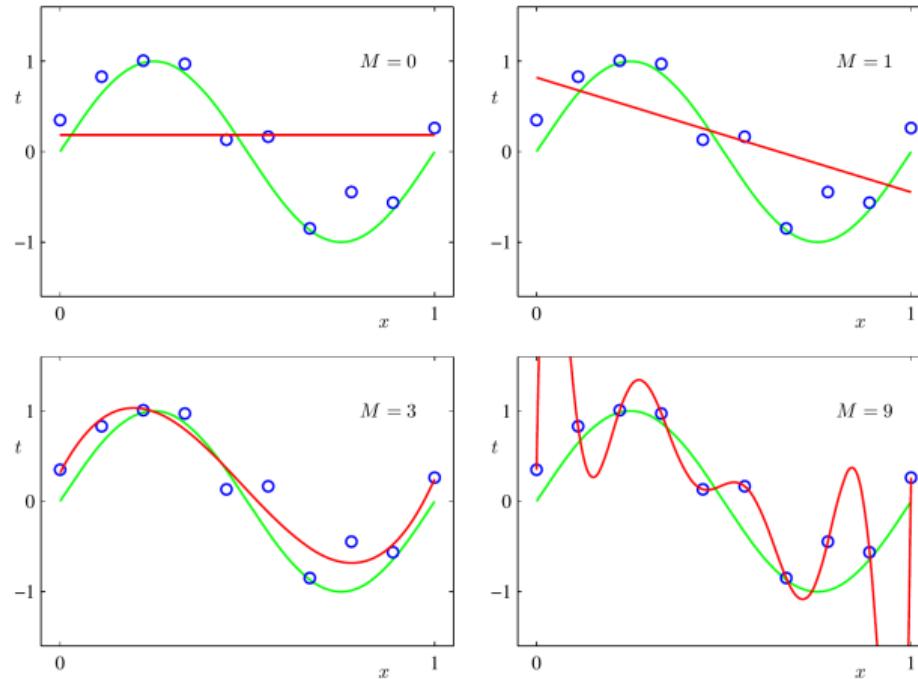
- ▶ How can this function be minimized wrt.  $\mathbf{w}$ ? Why is it easy? How many solutions?

# Polynomial Curve Fitting



The error function measures the displacement along the  $t$  dimension between the data points (blue) and the current model specified by the parameters  $\mathbf{w}$  (red).

# Polynomial Curve Fitting

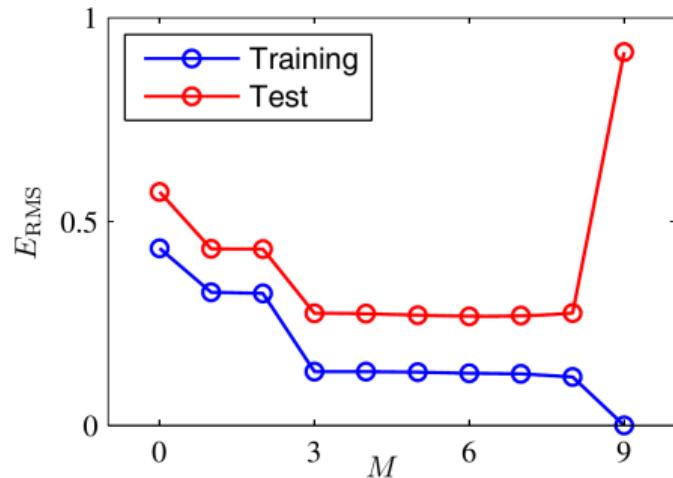


Plots of polynomials of various orders  $M$  (red) fitted to the dataset (blue). Underfitting ( $M = 0$ ) and overfitting ( $M = 9$ ) is observed. This is a model selection problem.

# Polynomial Curve Fitting

How can we measure generalization performance (overfitting or underfitting)?

- ▶ Consider performance on held-out test set from same distribution



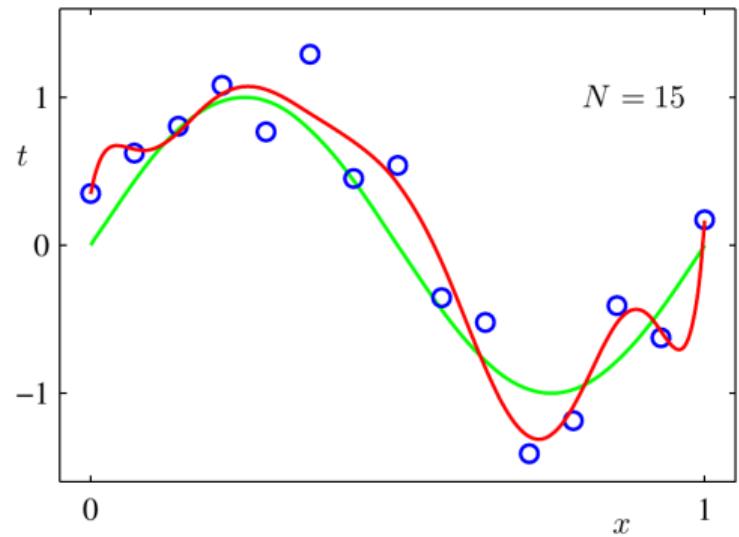
Root Mean Square (RMS) error evaluated on the training set (blue) and test set (red) for various  $M$ . Why do we obtain  $E_{RMS} = 0$  for  $M = 9$  on the training set?

# Polynomial Curve Fitting

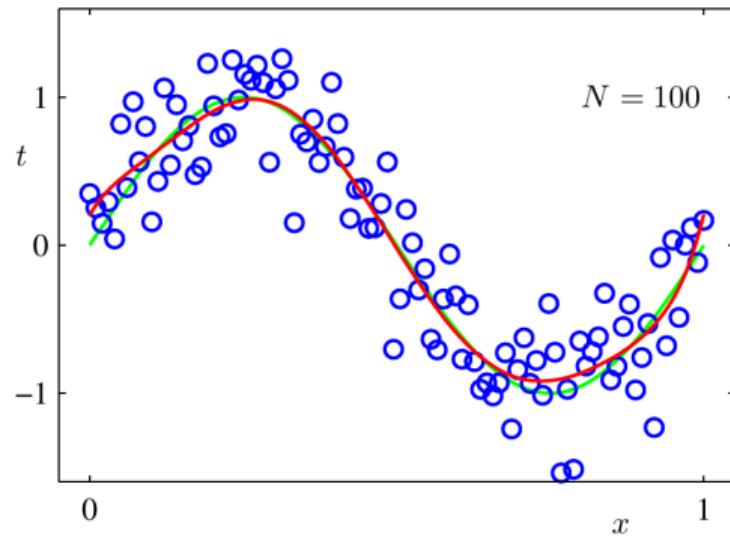
	$M = 0$	$M = 1$	$M = 6$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43

Polynomial coefficients  $\mathbf{w}$  for polynomials of various order fitted to the dataset.

# Polynomial Curve Fitting



$N = 15$



$N = 100$

Polynomial fits (red) obtained by minimizing the error function using polynomial order  $M = 9$  for a dataset of size 15 (left) and 100 (right). What is the effect of more data?

# Polynomial Curve Fitting

Model:

$$y(x, \mathbf{w}) = \sum_{j=0}^M w_j x^j$$

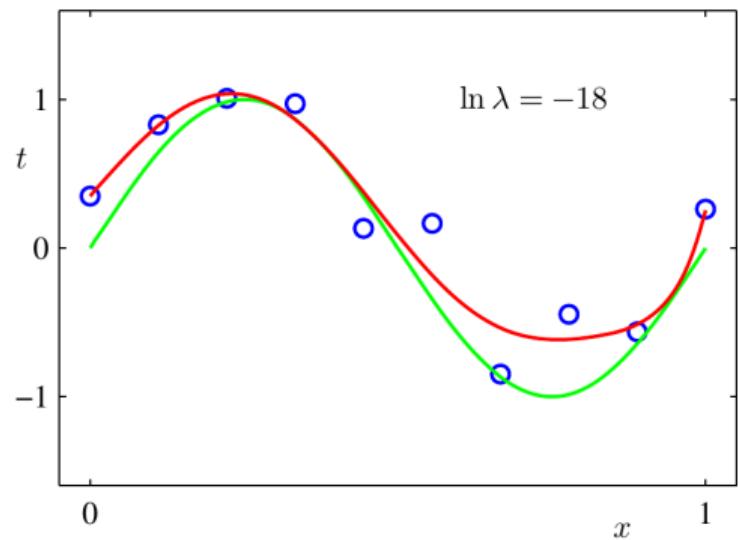
- ▶ Loss function:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2$$

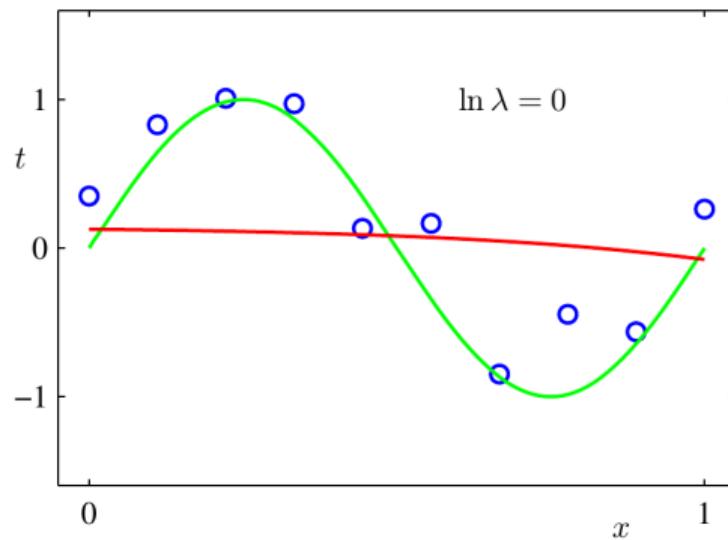
- ▶ How can we avoid overfitting without reducing  $K$ ?
- ▶ Regularizing the loss function (quadratic case = ridge regression):

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

# Polynomial Curve Fitting



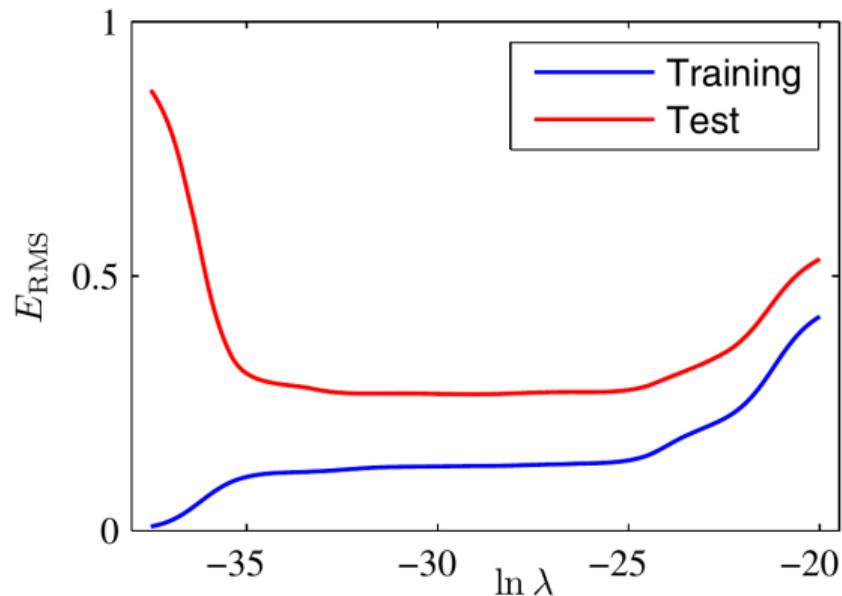
$$\ln \lambda = -18$$



$$\ln \lambda = 0$$

Regularized polynomial fits (red) with weak (left) and strong (right) regularization.

# Polynomial Curve Fitting



Root Mean Square (RMS) error evaluated on the training set (blue) and test set (red) for various regularization values  $\lambda$ . Which  $\lambda$  value would you pick?

# Non-Parametric Methods

# Parametric vs. Non-Parametric Methods

- ▶ Parametric Models
  - ▶ Select model with relatively small number of parameters
  - ▶ Fit model parameters to data (=training phase)
  - ▶ Generalize well if model is correctly specified
  - ▶ Do not require the data at inference time
  - ▶ Problematic if model is misspecified (e.g., Gaussian for multi-modal distributions)
- ▶ Non-Parametric Models
  - ▶ Do not require the definition of a parametric model
  - ▶ Instead, directly base their predictions on the dataset (often not even training phase)
  - ▶ Scale with the size of the dataset
  - ▶ Require more storage, often slower
  - ▶ However, often worse in terms of generalization (less assumptions specified)

# K-Nearest Neighbor Classification

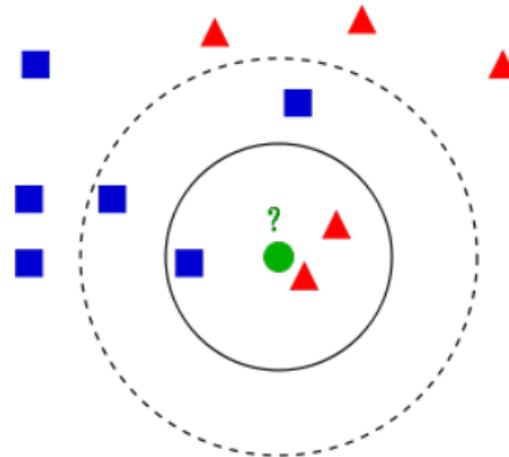
# K-NN Classification

## K-Nearest Neighbor Classification

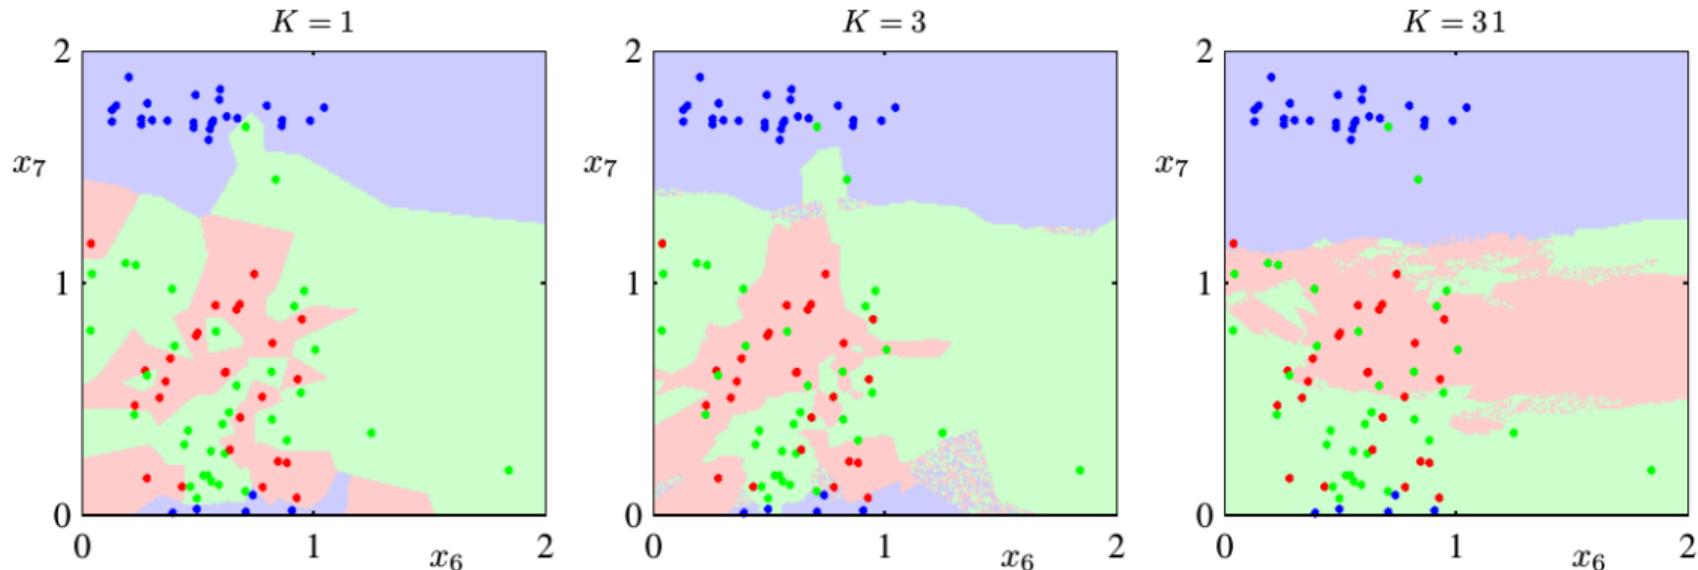
- ▶ Assign query point to the class most common among its K nearest neighbors

## Special case: Nearest Neighbor Classification

- ▶ Assign query point to class of nearest neighbor



# K-NN Classification



200 points from the oil dataset (Bishop and James, 1993). This plot shows 2 of 12 feature dimensions for 3 different classes varying  $K$ . What do you observe?

Questions?