

# Machine Learning in Graphics and Vision

Prof. Dr.-Ing. Andreas Geiger

Autonomous Vision Group  
MPI-IS / University of Tübingen

July 12, 2018



University of Tübingen  
MPI for Intelligent Systems  

---

Autonomous Vision Group



So far: Learning from very large Datasets



# Now: Integrate Prior Knowledge

## Goal

- ▶ Take probabilistic viewpoint
- ▶ Model dependency structure of problem
- ▶ Model relationships between entities parametrically
- ▶ Model complex phenomena by composing simple functions

## Pros

- ▶ Easy integration of prior knowledge
- ▶ Training with limited data
- ▶ Models often easily interpretable

## Cons

- ▶ Many phenomena are hard to model accurately

# Overview

## **Structured Prediction I**

- ▶ Graphical Models: Factor Graphs
- ▶ Inference: Belief Propagation

## **Structured Prediction II**

- ▶ Stereo & Optical Flow
- ▶ Multi-view Reconstruction

## **Structured Prediction III**

- ▶ Parameter Estimation
- ▶ Deep Structured Models

# Structured Prediction

# Non-Structured vs. Structured Prediction

## Classification / Regression:

$$f : \mathcal{X} \rightarrow \mathbb{N} \quad \text{or} \quad f : \mathcal{X} \rightarrow \mathbb{R}$$


- ▶ Inputs  $x \in \mathcal{X}$  can be any kind of objects
  - ▶ images, text, audio, sequence of amino acids, ...
- ▶ Output  $y \in \mathbb{N}/y \in \mathbb{R}$  is a discrete or real number
  - ▶ classification, regression, density estimation, ...

## Structured Output Learning:

$$f : \mathcal{X} \rightarrow \mathcal{Y}$$

- ▶ Inputs  $x \in \mathcal{X}$  can be any kind of objects
- ▶ Outputs  $y \in \mathcal{Y}$  are complex (structured) objects
  - ▶ images, text, parse trees, folds of a protein, computer programs, ...

# Supervised Learning




- **Learning:** Estimate parameters  $\theta$  from training dataset  $\{(x_i, y_i)\}_{i=1}^N$
- **Inference:** Make novel predictions  $f_{\theta}(\cdot)$

# Classification / Regression

Input

Model

Output



"Beach"


Classification / Regression:

- ▶ Output is encoded in a single one-dimensional variable
- ▶ Many problems are not of this form




# Structured Prediction


Input



Model



Output



Graphical models allow us to encode:

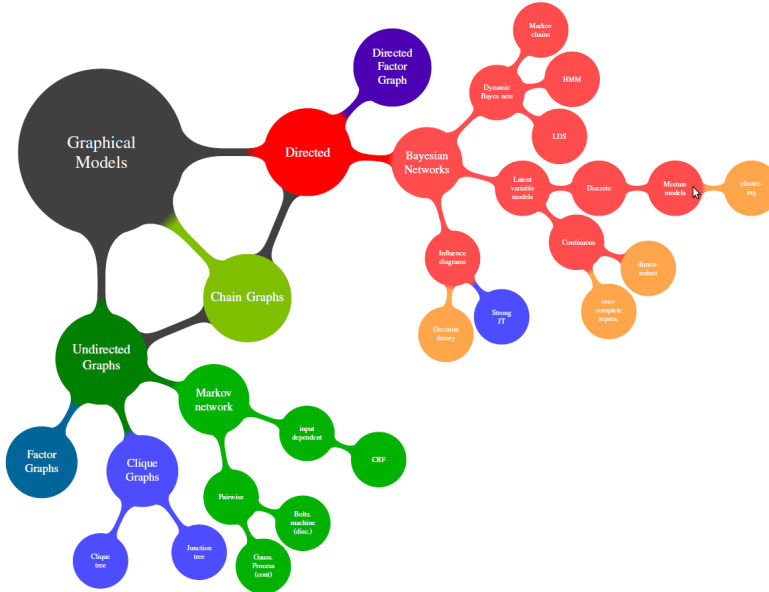
- ▶ The dependency structure of the problem
- ▶ Constraints between random variables

# Probabilistic Graphical Models

# Probabilistic Graphical Models

- ▶ "Graphical language" to represent probability distributions
- ▶ In particular:
  - ▶ Dependence/independence of variables
  - ▶ Causal relationships (directed models)
- ▶ Tool to specify prior knowledge (expert knowledge)
- ▶ Allow for efficient learning & inference
- ▶ Many algorithms available
- ▶ Many different types of graphical models
  - ▶ Choice depends on application

# Probabilistic Graphical Models



# Markov Networks

# Markov Network

## Potential

A **potential**  $\phi(x)$  is a non-negative function of the variable  $x$ . A **joint potential**  $\phi(x_1, \dots, x_D)$  is a non-negative function of the **set** of variables.


## Markov Network

For a set of variables  $\mathcal{X} = \{x_1, \dots, x_D\}$  a **Markov network** is defined as a product of potentials over the **maximal cliques**  $\mathcal{X}_c$  of the graph  $\mathcal{G}$

$$p(x_1, \dots, x_D) = \frac{1}{Z} \prod_{c=1}^C \phi_c(\mathcal{X}_c)$$

- ▶ Special case for clique size two: **Pairwise Markov network**
- ▶ If all potentials are strictly positive: **Gibbs distribution**

## Example




$$p(a, b, c) = \frac{1}{Z} \phi_1(a, c) \phi_2(b, c)$$

- ▶ Two maximal cliques of size two:  $\phi_1(a, c)$  and  $\phi_2(b, c)$
- ▶  $Z$  normalizes the distribution and is called **partition function**

$$Z = \sum_{a,b,c} \phi_1(a, c) \phi_2(b, c)$$

# Properties of Markov Networks




- Marginalizing over  $c$  makes  $a$  and  $b$  dependent
- Proof by checking

$$p(a, b) \neq p(a)p(b)$$



# Properties of Markov Networks



- ▶ Conditioning on  $c$  makes  $a$  and  $b$  independent
- ▶ Proof by checking

$$p(a, b \mid c) = p(a \mid c)p(b \mid c)$$

Now the general case ..


# Global Markov Property

## Separation

A subset  $\mathcal{S}$  separates  $\mathcal{A}$  from  $\mathcal{B}$  if every path from a member of  $\mathcal{A}$  to any member of  $\mathcal{B}$  passes through  $\mathcal{S}$ .

## Global Markov Property

For disjoint sets of variables  $(\mathcal{A}, \mathcal{B}, \mathcal{S})$  where  $\mathcal{S}$  separates  $\mathcal{A}$  from  $\mathcal{B}$ , we have  $\mathcal{A} \perp\!\!\!\perp \mathcal{B} \mid \mathcal{S}$




# Local Markov Property

## Local Markov Property

$$p(x \mid \mathcal{X} \setminus \{x\}) = p(x \mid ne(x))$$

- ▶ The set of neighboring nodes  $ne(x)$  is called **Markov blanket**
- ▶ This also holds for sets of variables

## Local Markov Property – Example



- ▶  $p(x_4 \mid x_1, x_2, x_3, x_4, x_5, x_6, x_7) = p(x_4 \mid x_2, x_3, x_5, x_6)$
- ▶ In other words  $x_4 \perp\!\!\!\perp \{x_1, x_7\} \mid \{x_2, x_3, x_5, x_6\}$
- ▶ And others ..

# Markov Random Field (MRF)

## Markov Random Field

A **Markov Random Field** is defined by a set of distributions  $p(x_i \mid ne(x_i))$  with respect to an undirected graph  $\mathcal{G}$  such that


$$p(x_i \mid x_{\setminus i}) = p(x_i \mid ne(x_i))$$

- Not every set of conditional distributions  $p(x_i \mid x_{\setminus i})$  yields a valid joint distribution

# Finding the factorization

- ▶ An undirected graph  $\mathcal{G}$  specifies a set of conditional independence statements
- ▶ Which factorization satisfies all possible independence assumptions?
- ▶ What is the most general factorization  $F$  that satisfies the independence assumptions of  $\mathcal{G}$ ?

## Finding the factorization



- Eliminate variable one by one
- Let's start with  $x_1$


$$p(x_1, \dots, x_7) = p(x_1 \mid x_2, x_3)p(x_2, \dots, x_7)$$

since

$$p(x_1 \mid x_2, \dots, x_7) = p(x_1 \mid x_2, x_3)$$



# Finding the factorization




- Graph specifies:

$$\begin{aligned} p(x_1, x_2, x_3 \mid x_4 \dots, x_7) &= p(x_1, x_2, x_3 \mid x_4) \\ \Rightarrow p(x_2, x_3 \mid x_4, \dots, x_7) &= p(x_2, x_3 \mid x_4) \end{aligned}$$

- Hence

$$p(x_1, \dots, x_7) = p(x_1 \mid x_2, x_3) p(x_2, x_3 \mid x_4) p(x_4, x_5, x_6, x_7)$$

## Finding the factorization




- We continue to find

$$\begin{aligned} p(x_1, \dots, x_7) &= p(x_1 \mid x_2, x_3) p(x_2, x_3 \mid x_4) \\ &\quad p(x_4 \mid x_5, x_6) p(x_5, x_6 \mid x_7) p(x_7) \end{aligned}$$

- Factorization into clique potentials (maximal cliques)

$$p(x_1, \dots, x_7) = \frac{1}{Z} \phi(x_1, x_2, x_3) \phi(x_2, x_3, x_4) \phi(x_4, x_5, x_6) \phi(x_5, x_6, x_7)$$

# Finding the factorization



- ▶ Markov conditions of  $\mathcal{G} \Rightarrow$  factorization  $F$  into cliques
- ▶ And conversely:  $F \Rightarrow \mathcal{G}$

# Hammersley-Clifford Theorem

## Hammersley-Clifford

Relationship Markov conditions on  $\mathcal{G} \Leftrightarrow$  Factorization  $F$  holds for any undirected graph provided that the potentials are positive

- ▶ Thus also loopy ones:  $x_1 - x_2 - x_3 - x_4 - x_1$
- ▶ Theorem says, distribution is of the form

$$p(x_1, x_2, x_3, x_4) = \frac{1}{Z} \phi_{12}(x_1, x_2) \phi_{23}(x_2, x_3) \phi_{34}(x_3, x_4) \phi_{41}(x_4, x_1)$$


# Factor Graphs

# Relationship Potentials to Graphs

- Consider this factorization into potential functions:

$$p(a, b, c) = \frac{1}{Z} \phi(a, b) \phi(b, c) \phi(c, a)$$

- What is the corresponding Markov network (graphical model)?




- and which other factorization is represented by this network?

$$p(a, b, c) = \frac{1}{Z} \phi(a, b, c)$$

- The factorization of the **potentials** is not uniquely specified by the graph

# Relationship Potentials to Graphs

- Now we introduce an extra node (a square) for each factor



- Left: Markov Network
- Middle: Factor graph representation of  $\phi(a, b, c)$
- Right: Factor graph representation of  $\phi(a, b)\phi(b, c)\phi(c, a)$
- Different factor graphs have same Markov network  $(b, c) \Rightarrow (a)$

# Factor Graph Definition

## Factor Graph

Given a function

$$f(x_1, \dots, x_n) = \prod_i f_i(\mathcal{X}_i)$$

the **factor graph (FG)** has a **square node** for each factor  $f_i(\mathcal{X}_i)$  and a **circle node** for each variable  $x_j$ . We typically specify this factorization up to a normalization constant


$$p(x_1, \dots, x_n) = \frac{1}{Z} \prod_i f_i(\mathcal{X}_i)$$

when representing a distribution  $p(\cdot)$ .



# Factor Graph: Example 1

- Question: which distribution ?



- Answer:


$$p(x) = \frac{1}{Z} f_a(x_1, x_2) f_b(x_1, x_2) f_c(x_2, x_3) f_d(x_3)$$

## Factor Graph: Example 2

- Question: Which factor graph?


$$p(x_1, x_2, x_3) = p(x_1) p(x_2) p(x_3 | x_1, x_2)$$

- Answer:



# Inference in Factor Graphs

# Inference in Chain Structured Factor Graphs




$$p(a, b, c, d) = \frac{1}{Z} f_1(a, b) f_2(b, c) f_3(c, d) f_4(d)$$

$$p(a) = \sum_{b, c, d} p(a, b, c, d) = ?$$

Computational Complexity?

$$\begin{aligned} p(a, b, c) &= \sum_d p(a, b, c, d) \\ &= \frac{1}{Z} f_1(a, b) f_2(b, c) \underbrace{\sum_d f_3(c, d) f_4(d)}_{\mu_{d \rightarrow c}(c)} \end{aligned}$$

# Inference in Chain Structured Factor Graphs




- ▶ Simply recurse further:

$$p(a) = \sum_b p(a, b) = \frac{1}{Z} \sum_b f_1(a, b) \mu_{c \rightarrow b}(b) = \frac{1}{Z} \mu_{b \rightarrow a}(a)$$

- ▶  $\mu_{m \rightarrow n}(n)$  carries the information beyond  $m$
- ▶ Computational complexity?
- ▶ We did not need the factors yet
- ▶ But we will see that making a distinction is helpful

# Inference in Tree Structured Factor Graphs

- Consider a branching graph:




with factors

$$f_1(a, b) f_2(b, c, d) f_3(c) f_4(d, e) f_5(d)$$


- How to find marginal  $p(a, b)$ ?

# Inference in Tree Structured Factor Graphs

- Idea: compute messages



# Inference in Tree Structured Factor Graphs




$$p(a, b) = \frac{1}{Z} f_1(a, b) \underbrace{\sum_{c, d, e} f_2(b, c, d) f_3(c) f_5(d) f_4(d, e)}_{\mu_{f_2 \rightarrow b}(b)}$$

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c, d} f_2(b, c, d) f_3(c) f_5(d) \sum_e f_4(d, e)$$




# Inference in Tree Structured Factor Graphs



$$p(a, b) = \frac{1}{Z} f_1(a, b) \underbrace{\sum_{c, d, e} f_2(b, c, d) f_3(c) f_5(d) f_4(d, e)}_{\mu_{f_2 \rightarrow b}(b)}$$

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c, d} f_2(b, c, d) \underbrace{f_3(c)}_{\mu_{c \rightarrow f_2}(c)} \underbrace{f_5(d) \sum_e f_4(d, e)}_{\mu_{d \rightarrow f_2}(d)}$$


# Inference in Tree Structured Factor Graphs



$$p(a, b) = \frac{1}{Z} f_1(a, b) \underbrace{\sum_{c, d, e} f_2(b, c, d) f_3(c) f_5(d) f_4(d, e)}_{\mu_{f_2 \rightarrow b}(b)}$$

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c, d} f_2(b, c, d) \mu_{c \rightarrow f_2}(c) \mu_{d \rightarrow f_2}(d)$$

# Factor-to-Variable Messages




- Here (repeated from last slide):

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c,d} f_2(b, c, d) \mu_{c \rightarrow f_2}(c) \mu_{d \rightarrow f_2}(d)$$

- More general:


$$\mu_{f \rightarrow x}(x) = \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \prod_{y \in \{ne(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$

# Variable-to-Factor Messages




$$\mu_{d \rightarrow f_2}(d) = f_5(d) \sum_e f_4(d, e)$$

# Variable-to-Factor Messages




$$\mu_{d \rightarrow f_2}(d) = \underbrace{f_5(d)}_{\mu_{f_5 \rightarrow d}(d)} \underbrace{\sum_e f_4(d, e)}_{\mu_{f_4 \rightarrow d}(d)}$$

# Variable-to-Factor Messages



$$\mu_{d \rightarrow f_2}(d) = \mu_{f_5 \rightarrow d}(d) \mu_{f_4 \rightarrow d}(d)$$

# Variable-to-Factor Messages



- Here (repeated from last slide):

$$\mu_{d \rightarrow f_2}(d) = \mu_{f_5 \rightarrow d}(d) \mu_{f_4 \rightarrow d}(d)$$

- General:

$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{ne(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$

# Comments

- ▶ Many subscripts, don't get confused :)
- ▶ Once computed, messages can be re-used
- ▶ Important observation: All marginals  $(p(c), p(d), p(c, d), \dots)$  can be written as a function of messages
- ▶ We need an algorithm to compute all messages
- ▶ For marginal inference: Sum-product algorithm



# Sum-Product Algorithm

# Sum-Product Algorithm – Overview

## **Belief Propagation:**


- ▶ Algorithm to compute all messages efficiently
- ▶ Assumes that the graph is singly-connected (chain, tree)

## **Algorithm:**

1. Initialization
2. Variable to Factor message
3. Factor to Variable message
4. Repeat until all messages have been calculated
5. Calculate the desired marginals from the messages


# 1. Initialization

- ▶ Messages from extremal node factors are initialized to factor
- ▶ Messages from extremal variable nodes can be set arbitrarily (e.g., to 1)




## 2. Variable-to-Factor Message

$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{ne(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$



### 3. Factor-to-Variable Message


$$\mu_{f \rightarrow x}(x) = \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \prod_{y \in \{ne(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$



- We sum over all states in the set of variables
- This explains the name for the algorithm (sum-product)
- Great, this is tractable now! Or not?

## 5. Calculate Marginals

$$p(x) \propto \prod_{f \in ne(x)} \mu_{f \rightarrow x}(x)$$



# Log Representation

- ▶ In large graphs, messages may become very small/big
- ▶ Work with log-messages instead  $\lambda = \log \mu$
- ▶ Variable-to-factor messages

$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{ne(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$

then becomes

$$\lambda_{x \rightarrow f}(x) = \sum_{g \in \{ne(x) \setminus f\}} \lambda_{g \rightarrow x}(x)$$

# Log Representation

- ▶ Work with log-messages instead  $\lambda = \log \mu$
- ▶ Factor-to-variable messages

$$\mu_{f \rightarrow x}(x) = \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \prod_{y \in \{ne(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$

then become

$$\lambda_{f \rightarrow x}(x) = \log \left( \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \exp \left[ \sum_{y \in \{ne(f) \setminus x\}} \lambda_{y \rightarrow f}(y) \right] \right)$$



# Max-Product Algorithm

## Finding the maximal state: Max-Product

- For a given distribution  $p(a, b, c, d)$  find the most likely state:

$$a^*, b^*, c^*, d^* = \operatorname{argmax}_{a, b, c, d} p(a, b, c, d)$$

- This is called the **Maximum-A-Posteriori (MAP)** solution
- Again use factorization structure to distribute maximisation to local computations
- Chain example:

$$p(a, b, c, d) = \frac{1}{Z} f_1(a, b) f_2(b, c) f_3(c, d)$$

## Example: Chain

$$\begin{aligned}\max_{a,b,c,d} p(a,b,c,d) &= \max_{a,b,c,d} f_1(a,b)f_2(b,c)f_3(c,d) \\&= \max_{a,b,c} f_1(a,b)f_2(b,c) \underbrace{\max_d f_3(c,d)}_{\mu_{d \rightarrow c}(c)} \\&= \max_{a,b} f_1(a,b) \underbrace{\max_c f_2(b,c)\mu_{d \rightarrow c}(c)}_{\mu_{c \rightarrow b}(b)} \\&= \max_a \underbrace{\max_b f_1(a,b)\mu_{c \rightarrow b}(b)}_{\mu_{b \rightarrow a}(a)} \\&= \max_a \mu_{b \rightarrow a}(a)\end{aligned}$$

► Is this what we wanted to compute in the beginning?

## Example: Chain

- Once messages are computed, find the optimal values:

$$a^* = \operatorname{argmax}_a \mu_{b \rightarrow a}(a)$$

$$b^* = \operatorname{argmax}_b f_1(a^*, b) \mu_{c \rightarrow b}(b)$$

$$c^* = \operatorname{argmax}_c f_2(b^*, c) \mu_{d \rightarrow c}(c)$$

$$d^* = \operatorname{argmax}_d f_3(c^*, d)$$

- This is called **backtracking** (dynamic programming)
- If maximum unique: MAP = max of "max-marginals"

# Max-Product Algorithm – Overview

## **Belief Propagation:**


- ▶ Algorithm to compute all messages efficiently
- ▶ Assumes that the graph is singly-connected (chain, tree)

## **Algorithm:**

1. Initialization
2. Variable to Factor message
3. Factor to Variable message
4. Repeat until all messages have been calculated
5. Calculate the desired MAP solution

# 1. Initialisation


- ▶ Messages from extremal node factors are initialized to factor
- ▶ Messages from extremal variable nodes can be set arbitrarily (e.g., to 1)



- ▶ Same as for sum-product

## 2. Variable to Factor message


$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{ne(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$



- Same as for sum-product

### 3. Factor to Variable message

$$\mu_{f \rightarrow x}(x) = \max_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \prod_{y \in \{ne(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$




- Different message than in sum-product
- This is now a max-product!



# Computing the Maximal State of a Variable

$$x^* = \operatorname{argmax}_x \prod_{f \in ne(x)} \mu_{f \rightarrow x}(x)$$



# Log Representation

- ▶ In large graphs, messages may become very small/big
- ▶ Work with log-messages instead  $\lambda = \log \mu$
- ▶ Note: This doesn't change the optimization problem since

$$\log \left( \max_x p(x) \right) = \max_x \log (p(x))$$

- ▶ Variable-to-factor messages

$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{ne(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$

then become

$$\lambda_{x \rightarrow f}(x) = \sum_{g \in \{ne(x) \setminus f\}} \lambda_{g \rightarrow x}(x)$$

# Log Representation

- ▶ Work with log-messages instead  $\lambda = \log \mu$
- ▶ Factor-to-variable messages

$$\mu_{f \rightarrow x}(x) = \max_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \prod_{y \in \{ne(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$

then become

$$\lambda_{f \rightarrow x}(x) = \max_{\mathcal{X}_f \setminus x} \left[ \log f(\mathcal{X}_f) + \sum_{y \in \{ne(f) \setminus x\}} \lambda_{y \rightarrow f}(y) \right]$$

- ▶ This algorithm is called the **max-sum algorithm**

What if the graph is not singly connected?

# Loopy Belief Propagation

# Loopy Belief Propagation

$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{ne(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$


$$\mu_{f \rightarrow x}(x) = \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \prod_{y \in \{ne(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$

- ▶ Messages are also well defined for loopy graphs!
- ▶ Simply apply them to loopy graphs as well
- ▶ We loose exactness ( $\Rightarrow$  approximate inference)
- ▶ Even no guarantee of convergence [Yedida et al. 2004]
- ▶ But often works surprisingly well in practice

# Loopy Belief Propagation

Which message passing schedule?

- ▶ Random or fixed order
- ▶ Popular choice:
  1. Factors  $\rightarrow$  variables
  2. Variables  $\rightarrow$  factors
  3. Repeat for  $N$  iterations
- ▶ Can be run in parallel as factor graph is bipartite:



# Summary



# Sum-Product Belief Propagation

- **Goal:** Compute **marginals** of distribution
- **Factor-to-variable messages:**

$$\lambda_{f \rightarrow x}(x) = \log \left( \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \exp \left\{ \sum_{y \in \{ne(f) \setminus x\}} \lambda_{y \rightarrow f}(y) \right\} \right) \quad (1)$$

- **Variable-to-factor messages:**

$$\lambda_{x \rightarrow f}(x) = \sum_{g \in \{ne(x) \setminus f\}} \lambda_{g \rightarrow x}(x) \quad (2)$$

- $\sum_{\mathcal{X}_f \setminus x}$  : Summation over all states of  $\mathcal{X}_f \setminus x$  (Eq. 1)
- $\sum_{y \in \{ne(f) \setminus x\}} / \sum_{g \in \{ne(x) \setminus f\}}$  : Summation over all incoming messages / factors
- To avoid large values, subtract mean from  $\lambda_{x \rightarrow f}(x)$  after message update (Eq. 2)

# Max-Product Belief Propagation

- **Goal:** Find **most likely state** (MAP state)
- **Factor-to-variable messages:**

$$\lambda_{f \rightarrow x}(x) = \max_{\mathcal{X}_f \setminus x} \left[ \log f(\mathcal{X}_f) + \sum_{y \in \{ne(f) \setminus x\}} \lambda_{y \rightarrow f}(y) \right] \quad (3)$$

- **Variable-to-factor messages:**

$$\lambda_{x \rightarrow f}(x) = \sum_{g \in \{ne(x) \setminus f\}} \lambda_{g \rightarrow x}(x) \quad (2)$$

- $\max_{\mathcal{X}_f \setminus x}$  : Maximization over all states of  $\mathcal{X}_f \setminus x$  (Eq. 3)
- $\sum_{y \in \{ne(f) \setminus x\}} / \sum_{g \in \{ne(x) \setminus f\}}$  : Summation over all incoming messages / factors
- To avoid large values, subtract mean from  $\lambda_{x \rightarrow f}(x)$  after message update (Eq. 2)

# Pairwise Case

Factor-to-variable messages simplify as follows.

Variable-to-factor messages don't simplify.

- ▶ **Sum-Product Belief Propagation:**

- ▶ Unary factor  $f(x)$ :

$$\lambda_{f \rightarrow x}(x) = \log f(x) \quad (1)$$

- ▶ Pairwise factor  $f(x, y)$ :

$$\lambda_{f \rightarrow x}(x) = \log \left( \sum_y f(x, y) \exp \{ \lambda_{y \rightarrow f}(y) \} \right) \quad (1)$$

# Pairwise Case

Factor-to-variable messages simplify as follows.

Variable-to-factor messages don't simplify.

- ▶ **Max-Product Belief Propagation:**

- ▶ Unary factor  $f(x)$ :

$$\lambda_{f \rightarrow x}(x) = \log f(x) \quad (3)$$

- ▶ Pairwise factor  $f(x, y)$ :

$$\lambda_{f \rightarrow x}(x) = \max_y [\log f(x, y) + \lambda_{y \rightarrow f}(y)] \quad (3)$$

# Readout

Read off **marginal** or **MAP state** at each variable:

- ▶ Similar to variable-to-factor messages
- ▶ However: summing over **all** incoming messages

$$p(x) = \exp\{\lambda(x)\} / \sum_x \exp\{\lambda(x)\} \quad (4)$$

$$\text{with } \lambda(x) = \sum_{g \in \{ne(x)\}} \lambda_{g \rightarrow x}(x)$$

$$x^* = \operatorname{argmax}_x \sum_{g \in \{ne(x)\}} \lambda_{g \rightarrow x}(x) \quad (5)$$

# Algorithm Overview

## Belief Propagation Algorithm


- ▶ Input: **variables** and **factors**
- ▶ Allocate all **messages**
- ▶ Initialize the **message** log values to 0 (=uniform distribution)
- ▶ For  $N$  iterations do
  - ▶ Update all **factor-to-variable** messages (Eq. 1 or Eq. 3)
  - ▶ Update all **variable-to-factor** messages (Eq. 2)
  - ▶ Normalize all **variable-to-factor** messages:
$$\mu_{x \rightarrow f}(x) \leftarrow \mu_{x \rightarrow f}(x) - \text{mean}(\mu_{x \rightarrow f}(x))$$
- ▶ Read off **marginal** or **MAP state** at each variable (Eq. 4 or Eq. 5)

Examples

## Example 1: Vehicle Localization




## Example 1: Vehicle Localization




- **Goal:** Estimate vehicle location at time  $t = 1, \dots, 10$
- **Variables:**  $\mathbf{x} = \{x_1, \dots, x_{10}\}$      $x_i \in \{1, 2, 3\}$
- **Observations:**  $\mathbf{o} = \{o_1, \dots, o_{10}\}$      $o_i \in \mathbb{R}^3$

## Example 1: Vehicle Localization



$$p_{\theta}(\mathbf{x}) = \frac{1}{Z} \prod_{i=1}^{10} f_i(x_i) \prod_{i=1}^9 g_{\theta}(x_i, x_{i+1})$$

## Example 1: Vehicle Localization




$$p_{\theta}(\mathbf{x}) = \frac{1}{Z} \prod_{i=1}^{10} f_i(x_i) \prod_{i=1}^9 g_{\theta}(x_i, x_{i+1})$$

### Unary Factors:

$$\blacktriangleright f_1(x_1) = \begin{bmatrix} 0.7 \\ 0.2 \\ 0.1 \end{bmatrix}, \quad f_2(x_2) = \begin{bmatrix} 0.7 \\ 0.1 \\ 0.2 \end{bmatrix}, \quad f_3(x_3) = \begin{bmatrix} 0.2 \\ 0.1 \\ 0.7 \end{bmatrix}, \quad \dots$$

## Example 1: Vehicle Localization



$$p_{\theta}(\mathbf{x}) = \frac{1}{Z} \prod_{i=1}^{10} f_i(x_i) \prod_{i=1}^9 g_{\theta}(x_i, x_{i+1})$$


### Pairwise Factors:

$$\triangleright g_{\theta}(x_i, x_{i+1}) = \begin{bmatrix} 0.8 & 0.2 & 0.0 \\ 0.2 & 0.6 & 0.2 \\ 0.0 & 0.2 & 0.8 \end{bmatrix}$$

► Learning Problem:

$$\theta^* = \operatorname{argmax}_{\theta} \prod_n p_{\theta}(\mathbf{x}_n | \mathbf{y}_n)$$

## Example 1: Vehicle Localization



- Maximum-A-Posteriori State:



$$\hat{x}_1, \dots, \hat{x}_{10} = \operatorname{argmax}_{x_1, \dots, x_{10}} p_{\theta}(x_1, \dots, x_{10})$$

- Marginal Distribution:

$$p(x_1) = \sum_{x_2} \sum_{x_3} \cdots \sum_{x_{10}} p_{\theta}(x_1, \dots, x_{10})$$



# Example 1: Vehicle Localization

## Observations



# Example 1: Vehicle Localization


**Marginal Distributions**




## Example 2: Image Denoising



## Example 2: Image Denoising




## Example 2: Image Denoising




## Example 2: Image Denoising


Can we recover the original image from a noisy observation?



- Variables:  $x_1, \dots, x_{100} \in \{0, 1\}$
- Unary potentials:  $\psi_1(x_1), \dots, \psi_{100}(x_{100})$
- $\psi_i(x_i) = [x_i = o_i]$  with observation  $o_i$
- Log representation:  $\psi_i(x_i) = \log f_i(x_i)$   
 $p(x) = \frac{1}{Z} \prod_i f_i(x_i) = \frac{1}{Z} \exp \{ \sum_i \psi_i(x_i) \}$




## Example 2: Image Denoising



- ▶ Let us look at the clean image again!
- ▶ What prior knowledge do we have about this image?

## Example 2: Image Denoising



- Log representation:

$$p(x) \propto \exp \left\{ \sum_{i=1}^{100} \psi_i(x_i) + \sum_{i \sim j} \psi_{ij}(x_i, x_j) \right\}$$

- Variables:  $x_1, \dots, x_{100} \in \{0, 1\}$
- Unaries:  $\psi_i(x_i) = [x_i = o_i]$  with observation  $o_i \in \{0, 1\}$
- Pairwise potential:  $\psi_{ij}(x_i, x_j) = \alpha \cdot [x_i = x_j]$
- Parameter  $\alpha$  controls strength of prior

Questions?