

# Machine Learning in Graphics and Vision

Prof. Dr.-Ing. Andreas Geiger

Autonomous Vision Group  
MPI-IS / University of Tübingen

July 26, 2018



University of Tübingen  
MPI for Intelligent Systems  

---

Autonomous Vision Group



# Lecture: Self-Driving Cars (WS 18/19)

## COMPUTER SCIENCE

Computational Systems Biology of Infections

Algorithmen der Bioinformatik

Kommunikationsnetze

**Lernbasierte Computer Vision**

Home

People

Research

Publications

Teaching

Lecture: Self-Driving Cars

Theses

Data Science & Analytics

Neuronale Informationsverarbeitung

Kognitive Systeme

Cognitive Modeling

Computergrafik

Symbolisches Rechnen

Human-Computer Interaction

## Lecture: Self-Driving Cars

### Content

This new course will give an introduction to self-driving cars. The course covers topics in perception, planning, control and end-to-end driving, amongst others.



### Overview

- SWS: 2 V + 2 Ü
- 6 ECTS
- Veranstaltungsnummer: INF

### News

- Please enroll in ILIAS at the beginning of the semester.

### Exercises

By continuous and active participation in the weekly exercises, students may obtain a 0.3 bonus on the final grade, when passing the exam. To qualify for this bonus, the student must successfully solve 60% of the assigned homework problems which will be determined by grading the submitted homework solutions.

### Lecturer

Prof. Dr. Andreas Geiger [✉](#)

### TAs

tbd

### Lecture Dates

Wintersemester 2018

### Exam Dates

→ tbd

# Overview

## **Structured Prediction I**

- ▶ Graphical Models: Factor Graphs
- ▶ Inference: Belief Propagation

## **Structured Prediction II**


- ▶ Stereo & Optical Flow
- ▶ Multi-view Reconstruction

## **Structured Prediction III**

- ▶ Parameter Estimation
- ▶ Deep Structured Models

# Materials

- ▶ Nowozin, Lampert: Structured Learning and Prediction in Computer Vision  
Foundations and Trends in Computer Graphics and Vision, Volume 6, Number 3-4
- ▶ <http://www.nowozin.net/sebastian/cvpr2012tutorial/>



Recap

# Factor Graph

## Factor Graph

Given a function

$$f(x_1, \dots, x_n) = \prod_i f_i(\mathcal{X}_i)$$


the **factor graph (FG)** has a **square node** for each factor  $f_i(\mathcal{X}_i)$  and a **circle node** for each variable  $x_j$ . We typically specify this factorization up to a normalization constant

$$p(x_1, \dots, x_n) = \frac{1}{Z} \prod_i f_i(\mathcal{X}_i)$$

when representing a distribution  $p(\cdot)$ .

# Example

- Factor Graph:




- Distribution:

$$p(x) = \frac{1}{Z} f_a(x_1, x_2) f_b(x_1, x_2) f_c(x_2, x_3) f_d(x_3)$$

# Marginal Inference

- Consider a branching graph:



with factors


$$f_1(a, b)f_2(b, c, d)f_3(c)f_4(d, e)f_5(d)$$

- How to find marginal  $p(a, b)$ ?




# Marginal Inference

- Idea: compute messages




# Marginal Inference



$$p(a, b) = \frac{1}{Z} f_1(a, b) \underbrace{\sum_{c, d, e} f_2(b, c, d) f_3(c) f_5(d) f_4(d, e)}_{\mu_{f_2 \rightarrow b}(b)}$$

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c, d} f_2(b, c, d) f_3(c) f_5(d) \sum_e f_4(d, e)$$


# Marginal Inference



$$p(a, b) = \frac{1}{Z} f_1(a, b) \underbrace{\sum_{c, d, e} f_2(b, c, d) f_3(c) f_5(d) f_4(d, e)}_{\mu_{f_2 \rightarrow b}(b)}$$

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c, d} f_2(b, c, d) \underbrace{f_3(c)}_{\mu_{c \rightarrow f_2}(c)} \underbrace{f_5(d) \sum_e f_4(d, e)}_{\mu_{d \rightarrow f_2}(d)}$$


# Marginal Inference



$$p(a, b) = \frac{1}{Z} f_1(a, b) \underbrace{\sum_{c, d, e} f_2(b, c, d) f_3(c) f_5(d) f_4(d, e)}_{\mu_{f_2 \rightarrow b}(b)}$$

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c, d} f_2(b, c, d) \mu_{c \rightarrow f_2}(c) \mu_{d \rightarrow f_2}(d)$$

# Marginal Inference




- Here (repeated from last slide):

$$\mu_{f_2 \rightarrow b}(b) = \sum_{c,d} f_2(b, c, d) \mu_{c \rightarrow f_2}(c) \mu_{d \rightarrow f_2}(d)$$

- More general:

$$\mu_{f \rightarrow x}(x) = \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \prod_{y \in \{ne(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$

# Marginal Inference



- Here (repeated from last slide):

$$\mu_{d \rightarrow f_2}(d) = \mu_{f_5 \rightarrow d}(d) \mu_{f_4 \rightarrow d}(d)$$

- General:

$$\mu_{x \rightarrow f}(x) = \prod_{g \in \{ne(x) \setminus f\}} \mu_{g \rightarrow x}(x)$$

# Log Representation

- ▶ Work with log-messages instead  $\lambda = \log \mu$
- ▶ Factor-to-variable messages

$$\mu_{f \rightarrow x}(x) = \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \prod_{y \in \{ne(f) \setminus x\}} \mu_{y \rightarrow f}(y)$$

then become

$$\lambda_{f \rightarrow x}(x) = \log \left( \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \exp \left[ \sum_{y \in \{ne(f) \setminus x\}} \lambda_{y \rightarrow f}(y) \right] \right)$$

# Sum-Product Belief Propagation

- **Goal:** Compute **marginals** of distribution
- **Factor-to-variable messages:**

$$\lambda_{f \rightarrow x}(x) = \log \left( \sum_{\mathcal{X}_f \setminus x} f(\mathcal{X}_f) \exp \left\{ \sum_{y \in \{ne(f) \setminus x\}} \lambda_{y \rightarrow f}(y) \right\} \right) \quad (1)$$

- **Variable-to-factor messages:**

$$\lambda_{x \rightarrow f}(x) = \sum_{g \in \{ne(x) \setminus f\}} \lambda_{g \rightarrow x}(x) \quad (2)$$

- $\sum_{\mathcal{X}_f \setminus x}$  : Summation over all states of  $\mathcal{X}_f \setminus x$  (Eq. 1)
- $\sum_{y \in \{ne(f) \setminus x\}} / \sum_{g \in \{ne(x) \setminus f\}}$  : Summation over all incoming messages / factors
- To avoid large values, subtract mean from  $\lambda_{x \rightarrow f}(x)$  after message update (Eq. 2)



# Max-Product Belief Propagation

- **Goal:** Find **most likely state** (MAP state)
- **Factor-to-variable messages:**

$$\lambda_{f \rightarrow x}(x) = \max_{\mathcal{X}_f \setminus x} \left[ \log f(\mathcal{X}_f) + \sum_{y \in \{ne(f) \setminus x\}} \lambda_{y \rightarrow f}(y) \right] \quad (3)$$

- **Variable-to-factor messages:**

$$\lambda_{x \rightarrow f}(x) = \sum_{g \in \{ne(x) \setminus f\}} \lambda_{g \rightarrow x}(x) \quad (2)$$

- $\max_{\mathcal{X}_f \setminus x}$  : Maximization over all states of  $\mathcal{X}_f \setminus x$  (Eq. 3)
- $\sum_{y \in \{ne(f) \setminus x\}} / \sum_{g \in \{ne(x) \setminus f\}}$  : Summation over all incoming messages / factors
- To avoid large values, subtract mean from  $\lambda_{x \rightarrow f}(x)$  after message update (Eq. 2)

# Readout

Read off **marginal** or **MAP state** at each variable:

- ▶ Similar to variable-to-factor messages
- ▶ However: summing over **all** incoming messages

$$p(x) = \exp\{\lambda(x)\} / \sum_x \exp\{\lambda(x)\} \quad (4)$$

$$\text{with } \lambda(x) = \sum_{g \in \{ne(x)\}} \lambda_{g \rightarrow x}(x)$$

$$x^* = \operatorname{argmax}_x \sum_{g \in \{ne(x)\}} \lambda_{g \rightarrow x}(x) \quad (5)$$


# Inference Algorithm Overview

## Belief Propagation Algorithm

- ▶ Input: **variables** and **factors**
- ▶ Allocate all **messages**
- ▶ Initialize the **message** log values to 0 (=uniform distribution)
- ▶ For  $N$  iterations do
  - ▶ Update all **factor-to-variable** messages (Eq. 1 or Eq. 3)
  - ▶ Update all **variable-to-factor** messages (Eq. 2)
  - ▶ Normalize all **variable-to-factor** messages:
$$\mu_{x \rightarrow f}(x) \leftarrow \mu_{x \rightarrow f}(x) - \text{mean}(\mu_{x \rightarrow f}(x))$$
- ▶ Read off **marginal** or **MAP state** at each variable (Eq. 4 or Eq. 5)

# Log-Linear Models


# Example 1



$$\begin{aligned} p(x) &= \frac{1}{Z} f_a(x_1, x_2) f_b(x_1, x_2) f_c(x_2, x_3) f_d(x_3) \\ &= \frac{1}{Z} \exp \{ \log (f_a(x_1, x_2) f_b(x_1, x_2) f_c(x_2, x_3) f_d(x_3)) \} \\ &= \frac{1}{Z} \exp \{ \log f_a(x_1, x_2) + \log f_b(x_1, x_2) + \log f_c(x_2, x_3) + \log f_d(x_3) \} \\ &= \frac{1}{Z} \exp \{ \psi_a(x_1, x_2) + \psi_b(x_1, x_2) + \psi_c(x_2, x_3) + \psi_d(x_3) \} \end{aligned}$$

► Log factors:  $\psi_i(x_i) = \log f_i(x_i)$

## Example 2: Image Denoising




- Factor representation:

$$p(x) \propto \prod_{i=1}^{100} f_i(x_i) \prod_{i \sim j} f_{ij}(x_i, x_j)$$

- Variables:  $x_1, \dots, x_{100} \in \{0, 1\}$
- Unary factors:  $f_i(x_i)$
- Pairwise factors:  $f_{ij}(x_i, x_j)$

## Example 2: Image Denoising




- Factor representation:

$$p(x) \propto \exp \left\{ \sum_{i=1}^{100} \log f_i(x_i) \sum_{i \sim j} \log f_{ij}(x_i, x_j) \right\}$$

- Variables:  $x_1, \dots, x_{100} \in \{0, 1\}$
- Unary factors:  $f_i(x_i)$
- Pairwise factors:  $f_{ij}(x_i, x_j)$

## Example 2: Image Denoising



- Log-linear representation:

$$p(x) \propto \exp \left\{ \sum_{i=1}^{100} \psi_i(x_i) + \alpha \sum_{i \sim j} \psi_{ij}(x_i, x_j) \right\}$$

- Factors  $f(\cdot)$  become potentials  $\psi(\cdot)$  in log-representation
- Unary potentials:  $\psi_i(x_i) = [x_i = o_i]$  with observation  $o_i \in \{0, 1\}$
- Pairwise potentials:  $\psi_{ij}(x_i, x_j) = [x_i = x_j]$
- Parameter  $\alpha$  controls strength of prior – how to choose  $\alpha$ ? Learn from data!



# Parameter Estimation

# Factor Graph: Inference vs. Learning

$$p(x_1, \dots, x_{100}) = \frac{1}{Z} \exp \left\{ \sum_{i=1}^{100} \psi_i(x_i) + \alpha \sum_{i \sim j} \psi_{ij}(x_i, x_j) \right\}$$

- ▶ So far: Inference
  - ▶ Marginal distributions:  $p(x_i) = \sum_{x \setminus x_i} p(x_1, \dots, x_{100})$
  - ▶ MAP solution:  $x_1^*, \dots, x_{100}^* = \operatorname{argmax}_{x_1, \dots, x_{100}} p(x_1, \dots, x_{100})$
- ▶ Now: Learning
  - ▶ Estimate parameters (here:  $\alpha$ ) from dataset

# Conditional Random Fields

## Markov Random Field:

$$p(x) = \frac{1}{Z} \exp \left\{ \sum_{i=1}^{100} \psi_i(x_i) + \alpha \sum_{i \sim j} \psi_{ij}(x_i, x_j) \right\}$$

- Reason about output variables  $x \in \mathcal{X}$  given one particular model instantiation

## Structured Output Learning:

$$f : \mathcal{X} \rightarrow \mathcal{Y}$$

- Inputs  $x \in \mathcal{X}$  can be any kind of objects
- Outputs  $y \in \mathcal{Y}$  are complex (structured) objects
  - images, text, parse trees, folds of a protein, computer programs, ...

# Conditional Random Fields

## Markov Random Field:

$$p(x) = \frac{1}{Z} \exp \left\{ \sum_{i=1}^{100} \psi_i(x_i) + \alpha \sum_{i \sim j} \psi_{ij}(x_i, x_j) \right\}$$

- Reason about output variables  $x \in \mathcal{X}$  given one particular model instantiation

## Conditional Random Field:

$$p(y|x, w) = \frac{1}{Z} \exp \left\{ \sum_{i=1}^{100} \psi_i(x, y_i) + \alpha \sum_{i \sim j} \psi_{ij}(x, y_i, y_j) \right\}$$

- Make conditioning of variables  $y$  on data  $x$  and parameters  $w$  explicit (here  $w = \alpha$ )
- **MRF notation: outputs  $x \in \mathcal{X} \Rightarrow$  CRF notation: inputs  $x \in \mathcal{X}$ , outputs  $y \in \mathcal{Y}$**
- Learning: Estimate  $w$  from dataset  $\mathcal{D} = \{(x^1, y^1), \dots, (x^N, y^N)\}$

# Conditional Random Fields

## Conditional Random Field – General Form:

$$p(y|x, w) = \frac{1}{Z(x, w)} \exp \{ \langle w, \psi(x, y) \rangle \}$$

- ▶ Feature function:  $\psi(x, y) : \mathcal{X} \times \mathbb{R}^M \rightarrow \mathbb{R}^K$
- ▶ Parameter vector:  $w \in \mathbb{R}^K$  ( $M$ : num. output nodes,  $K$ : dim. feat. space)
- ▶ Partition function:  $Z(x, w) = \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x, y) \rangle \}$
- ▶ Note: much more flexible than just a single  $\alpha$  parameter!
- ▶ Learning: Estimate  $w$  from dataset  $\mathcal{D} = \{(x^1, y^1), \dots, (x^N, y^N)\}$
- ▶ GM specifies decomposition of  $\psi$  into potentials (=log factors)  $\psi_f$ :

$$\psi(x, y) = (\psi_f(x, y_f), \dots, \psi_{|\mathcal{F}|}(x, y_{|\mathcal{F}|}))$$

# Parameter Estimation

**Goal:** Maximize likelihood of outputs  $x$  conditioned on inputs  $y$  wrt.  $w$ :

$$w^* = \operatorname{argmax}_{w \in \mathbb{R}^K} p(y|x, w) \quad \text{with} \quad p(y|x, w) = \prod_{n=1}^N p(y^n|x^n, w)$$

This is equivalent to minimizing the negative conditional log-likelihood:

$$w^* = \operatorname{argmin}_{w \in \mathbb{R}^K} \mathcal{L}(w) \quad \text{with} \quad \mathcal{L}(w) = - \sum_{n=1}^N \log p(y^n|x^n, w)$$

# Parameter Estimation

**Goal:** Minimize negative conditional log-likelihood  $\mathcal{L}(w)$

$$w^* = \operatorname{argmin}_{w \in \mathbb{R}^K} \mathcal{L}(w)$$

$$\begin{aligned} \mathcal{L}(w) &= - \sum_{n=1}^N \log p(y^n | x^n, w) \\ &= - \sum_{n=1}^N \left[ \log \frac{1}{Z(x^n, w)} \exp \{ \langle w, \psi(x^n, y^n) \rangle \} \right] \\ &= - \sum_{n=1}^N [-\log Z(x^n, w) + \langle w, \psi(x^n, y^n) \rangle] \\ &= - \sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \end{aligned}$$


# Optimization

## Gradient Descent:

- ▶ Pick step size  $\eta$ , tolerance  $\epsilon$
- ▶ Initialize  $w^0 = 0$
- ▶ Repeat until  $\|v\| < \epsilon$ 
  - ▶  $v = \nabla_w \mathcal{L}(w)$
  - ▶  $\eta = \operatorname{argmin}_{\eta \in \mathbb{R}} \mathcal{L}(w^t - \eta v)$
  - ▶  $w^{t+1} = w^t - \eta v$

## Alternatives:

- ▶ Conjugate gradient
- ▶ L-BFGS
- ▶ All require gradient!





# Gradient of Negative Conditional Log-Likelihood

$$\begin{aligned}\mathcal{L}(w) &= -\sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \\ \nabla_w \mathcal{L}(w) &= -\sum_{n=1}^N \left[ \psi(x^n, y^n) - \frac{\sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \psi(x^n, y)}{\sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \}} \right] \\ &= -\sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} \frac{\exp \{ \langle w, \psi(x^n, y) \rangle \}}{\sum_{y' \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y') \rangle \}} \psi(x^n, y) \right] \\ &= -\sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) \right] \\ &= -\sum_{n=1}^N [\psi(x^n, y^n) - \mathbb{E}_{y \sim p(y|x^n, w)} \psi(x^n, y)]\end{aligned}$$

# Gradient of Negative Conditional Log-Likelihood

$$\nabla_w \mathcal{L}(w) = - \sum_{n=1}^N [\psi(x^n, y^n) - \mathbb{E}_{y \sim p(y|x^n, w)} \psi(x^n, y)]$$

When is  $\mathcal{L}(w)$  minimal?

$$\mathbb{E}_{y \sim p(y|x^n, w)} \psi(x^n, y) = \psi(x^n, y^n) \Rightarrow \nabla_w \mathcal{L}(w) = 0$$

- Interpretation: we aim at **expectation matching**:  $\mathbb{E}_{y \sim p} \psi(x, y) = \psi(x, y^{\text{obs}})$ , but discriminatively: only for  $x \in \{x^1, \dots, x^N\}$

Note:

- Hessian is pos. definite  $\Rightarrow \mathcal{L}(w)$  convex  $\Rightarrow \nabla_w \mathcal{L}(w) = 0$  implies global optimum!
- Only true as  $p(y|x, w)$  is log-linear in  $w \in \mathbb{R}^K$

# Computational Complexity

**Tasks** for gradient descent with line search:

- ▶ Evaluate  $\mathcal{L}(w)$
- ▶ Compute  $v = \nabla_w \mathcal{L}(w)$

$$\begin{aligned}\mathcal{L}(w) &= -\sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \\ \nabla_w \mathcal{L}(w) &= -\sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) \right]\end{aligned}$$

**Problem:**  $\mathcal{Y}$  is typically very (exponentially) large!

- ▶ Binary image segmentation:  $|\mathcal{Y}| = 2^{640 \times 480} \approx 10^{92475}$
- ▶ We must use the structure in  $\mathcal{Y}$ , or we are lost!

# Computational Complexity

$$\begin{aligned}\mathcal{L}(w) &= -\sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \\ \nabla_w \mathcal{L}(w) &= -\sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) \right]\end{aligned}$$

Computational complexity:  $O(ND^M K)$

- ▶  $N$ : number of samples in dataset ( $\approx 100$  to  $1,000,000$ )
- ▶  $M$ : number of output nodes ( $\approx 100$  to  $1,000,000$ )
- ▶  $D$ : maximal number of labels per output node ( $\approx 2$  to  $100$ )
- ▶  $K$ : dimensionality of feature space

# Computational Complexity

$$\begin{aligned}\mathcal{L}(w) &= -\sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \\ \nabla_w \mathcal{L}(w) &= -\sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) \right]\end{aligned}$$

Computational complexity:  $O(N \textcolor{red}{D}^M K)$

- ▶  $N$ : number of samples in dataset ( $\approx 100$  to  $1,000,000$ )
- ▶  $\textcolor{red}{M}$ : number of output nodes ( $\approx 100$  to  $1,000,000$ )
- ▶  $\textcolor{red}{D}$ : maximal number of labels per output node ( $\approx 2$  to  $100$ )
- ▶  $K$ : dimensionality of feature space

# Probabilistic Inference to the Rescue

Remember: in a graphical model, features decompose as follows

$$\psi(x, y) = (\psi_f(x, y_f), \dots, \psi_{|\mathcal{F}|}(x, y_{|\mathcal{F}|}))$$

Thus:

$$\begin{aligned} \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} &= \sum_{y \in \mathcal{Y}} \exp \left\{ \sum_{f \in \mathcal{F}} \langle w_f, \psi_f(x^n, y_f) \rangle \right\} \\ &= \sum_{y \in \mathcal{Y}} \prod_{f \in \mathcal{F}} \underbrace{\exp \{ \langle w_f, \psi_f(x^n, y_f) \rangle \}}_{\text{factor } f(\cdot)} \end{aligned}$$

- Can be efficiently calculated/approximated using message passing (run unnormalized sum-product BP, sum over any unnormalized marginal)

# Probabilistic Inference to the Rescue

Furthermore:

$$\begin{aligned}\sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) &= \mathbb{E}_{y \sim p(y|x^n, w)} \sum_{f \in \mathcal{F}} \psi_f(x^n, y_f) \\ &= \sum_{f \in \mathcal{F}} \mathbb{E}_{y \sim p(y|x^n, w)} \psi_f(x^n, y_f) \\ &= \sum_{f \in \mathcal{F}} \mathbb{E}_{y_f \sim p(y_f|x^n, w)} \psi_f(x^n, y_f) \\ &= \underbrace{\sum_{f \in \mathcal{F}}}_{|\mathcal{F}|} \underbrace{\sum_{y_f \in \mathcal{Y}_f}}_{D^F} \underbrace{p(y_f|x^n, w)}_{\text{marginal}} \psi_f(x^n, y_f)\end{aligned}$$

- ▶  $|\mathcal{F}|$ : number of factors,  $D$ : max. number of labels,  $F$ : order of largest factor
- ▶ Marginals can be calculated efficiently in polynomial time (e.g., with BP)

# Computational Complexity

$$\begin{aligned}\mathcal{L}(w) &= -\sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \\ \nabla_w \mathcal{L}(w) &= -\sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) \right]\end{aligned}$$

Computational complexity:  $\mathcal{O}(N \cancel{D^M} K) \rightarrow \mathcal{O}(N |\mathcal{F}| D^F K)$

- ▶  $N$ : number of samples in dataset ( $\approx 100$  to  $1,000,000$ )
- ▶  $M$ : number of output nodes ( $\approx 100$  to  $1,000,000$ )
- ▶  $D$ : maximal number of labels per output node ( $\approx 2$  to  $100$ )
- ▶  $K$ : dim. of feature space,  $|\mathcal{F}|$ : number of factors,  $F$ : order of largest factor



# Computational Complexity

$$\begin{aligned}\mathcal{L}(w) &= - \sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \\ \nabla_w \mathcal{L}(w) &= - \sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) \right]\end{aligned}$$

Computational complexity:  $O(N|\mathcal{F}|D^F K)$

- ▶  $N$ : number of samples in dataset ( $\approx 100$  to  $1,000,000$ )
- ▶  $M$ : number of output nodes ( $\approx 100$  to  $1,000,000$ )
- ▶  $D$ : maximal number of labels per output node ( $\approx 2$  to  $100$ )
- ▶  $K$ : dim. of feature space,  $|\mathcal{F}|$ : number of factors,  $F$ : order of largest factor

# Computational Complexity

## Large Datasets

- ▶ Processing all  $N$  training samples for one gradient update is slow

## How can we estimate parameters efficiently?

- ▶ Simplify model to make gradient updates faster  $\Rightarrow$  results get worse
- ▶ Train model on subsampled dataset  $\Rightarrow$  ignores information
- ▶ Parallelize across CPUs/GPUs  $\Rightarrow$  bottlenecks, doesn't save computation
- ▶ Stochastic gradient descent

# Stochastic Gradient Descent (SGD)

Stochastic Gradient Descent:

- ▶ Keep maximizing  $p(w|\mathcal{D})$
- ▶ In each gradient step:
  - ▶ Create random subset  $\mathcal{D}' \subset \mathcal{D}$  (typically  $1 \leq |\mathcal{D}'| \leq 64$ )
  - ▶ Follow approximate gradient:

$$\nabla_w \approx - \sum_{(x^n, y^n) \in \mathcal{D}'} [\psi(x^n, y^n) - \mathbb{E}_{y \sim p(y|x^n, w)} \psi(x^n, y)]$$

Comments:

- ▶ Line search no longer possible  $\Rightarrow$  extra hyper-parameter  $\eta$
- ▶ SGD converges to  $\operatorname{argmin}_w \mathcal{L}(w)$ ! (if  $\eta$  chosen right)
- ▶ SGD needs more iterations, but each one is faster
- ▶ See also: Bottou & Bousquet: The Tradeoffs of Large Scale Learning, NIPS 2007

# Computational Complexity

$$\begin{aligned}\mathcal{L}(w) &= - \sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \\ \nabla_w \mathcal{L}(w) &= - \sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) \right]\end{aligned}$$

Computational complexity:  $O(N|\mathcal{F}|D^F K)$

- ▶  $N$ : number of samples in dataset ( $\approx 100$  to  $1,000,000$ )
- ▶  $M$ : number of output nodes ( $\approx 100$  to  $1,000,000$ )
- ▶  $D$ : maximal number of labels per output node ( $\approx 2$  to  $100$ )
- ▶  $K$ : dim. of feature space,  $|\mathcal{F}|$ : number of factors,  $F$ : order of largest factor

# Computational Complexity

$$\begin{aligned}\mathcal{L}(w) &= -\sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right] \\ \nabla_w \mathcal{L}(w) &= -\sum_{n=1}^N \left[ \psi(x^n, y^n) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \psi(x^n, y) \right]\end{aligned}$$

Computational complexity:  $O(N|\mathcal{F}|D^F K)$

- ▶  $N$ : number of samples in dataset ( $\approx 100$  to  $1,000,000$ )
- ▶  $M$ : number of output nodes ( $\approx 100$  to  $1,000,000$ )
- ▶  $D$ : maximal number of labels per output node ( $\approx 2$  to  $100$ )
- ▶  $K$ : dim. of feature space,  $|\mathcal{F}|$ : number of factors,  $F$ : order of largest factor


# Feature Functions

## Semantic Segmentation:

- ▶  $\psi_i(x, y_i) \in \mathbb{R}^{\approx 1000}$ : local image features (e.g., bag of words)  
→  $\langle w_i, \psi_i(x, y_i) \rangle$ : local classifier (like logistic regression)
- ▶  $\psi_{i,j}(y_i, y_j) = [y_i = y_j] \in \mathbb{R}^1$ : test for same label  
→  $\langle w_{i,j}, \psi_{i,j}(y_i, y_j) \rangle$ : penalizer for label changes (if  $w_{i,j} > 0$ )
- ▶ combined:  $\operatorname{argmax}_y p(y|x, w)$  is smoothed version of local cues



original



local classification




local + smoothness

# Feature Functions

## Handwriting Recognition:

- ▶  $\psi_i(x, y_i) \in \mathbb{R}^{\approx 1000}$ : image representation (e.g., pixels, gradients)  
→  $\langle w_i, \psi_i(x, y_i) \rangle$ : local classifier for letters
- ▶  $\psi_{i,j}(y_i, y_j) \in \mathbb{R}^{26 \times 26}$ : letter/letter indicator  
→  $\langle w_{i,j}, \psi_{i,j}(y_i, y_j) \rangle$ : encourage/suppress letter combinations
- ▶ Combined:  $\operatorname{argmax}_y p(y|x, w)$  is “corrected” version of local cues



# Feature Functions

## Pose Estimation:

- ▶  $\psi_i(x, y_i) \in \mathbb{R}^{\approx 1000}$ : image representation (e.g., HoG)  
→  $\langle w_i, \psi_i(x, y_i) \rangle$ : local confidence map
- ▶  $\psi_{i,j}(y_i, y_j) = \text{fit}(y_i, y_j) \in \mathbb{R}^1$ : test for geometric fit  
→  $\langle w_{i,j}, \psi_{i,j}(y_i, y_j) \rangle$ : penalizer for unrealistic poses
- ▶ Combined:  $\text{argmax}_y p(y|x, w)$  is sanitized version of local cues



original



local classification



local + geometry



# Feature Functions

Typical feature functions for CRFs in computer vision:

- ▶ Unary terms  $\psi_i(x, y_i)$ : local representation, high-dimensional  
→  $\langle w_i, \psi_i(x, y_i) \rangle$ : local classifier
- ▶ Pairwise terms  $\psi_{i,j}(y_i, y_j)$ : prior knowledge, typically low-dimensional  
→  $\langle w_{i,j}, \psi_{i,j}(y_i, y_j) \rangle$ : penalize inconsistencies
- ▶ Pairwise terms sometimes also depend on  $x$ :  $\psi_{i,j}(x, y_i, y_j)$

Learning adjusts parameters:

- ▶ Unary weights  $w_i$ : learn local linear classifiers
- ▶ Pairwise weights  $w_{i,j}$ : learn importance of smoothing/penalization
- ▶  $\operatorname{argmax}_y p(y|x, w)$  is cleaned up version of local prediction

# Piece-wise Training

Sometimes, training the entire model at once is not easy:

- ▶ If terms depend on parameters in non-linear fashion
- ▶ If feature representations are high-dimensional, learning can be very slow

Alternative: Piece-wise Training

- ▶ Pre-train classifiers  $p(y_i|x)$ ; set  $\psi_i(x, y_i) = \log p(y_i|x) \in \mathbb{R}$
- ▶ Learn one-dimensional weight per classifier:  $\langle w_i, \psi_i(x, y_i) \rangle$

Advantage:

- ▶ Lower dimensional feature vector during training/inference  $\rightarrow$  faster
- ▶  $\log p(y_i|x)$  can be stronger classifiers, e.g., non-linear SVMs, CNNs, ..

Disadvantage

- ▶ If local classifiers are bad, CRF training cannot fix this

# Summary

Given:

- ▶ Training set  $\mathcal{D} = \{(x^1, y^1), \dots, (x^N, y^N)\}$  with  $(x^n, y^n) \stackrel{\text{i.i.d.}}{\sim} d(x, y)$
- ▶ Feature function:  $\psi(x, y) : \mathcal{X} \times \mathbb{R}^M \rightarrow \mathbb{R}^K$

Task:

- ▶ Find parameter vector  $w$  such that  $p(y|x, w) = \frac{1}{Z(x, w)} \exp \{ \langle w, \psi(x, y) \rangle \} \approx d(y|x)$

Minimize negative conditional log-likelihood:

$$\mathcal{L}(w) = - \sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right]$$

- ▶ Convex optimization problem  $\rightarrow$  gradient descent works
- ▶ Training needs repeated runs of probabilistic inference, faster is better

# Summary

Gradient of negative conditional log-likelihood:

$$\mathcal{L}(w) = - \sum_{n=1}^N \left[ \langle w, \psi(x^n, y^n) \rangle - \log \sum_{y \in \mathcal{Y}} \exp \{ \langle w, \psi(x^n, y) \rangle \} \right]$$

Problem	Solution	Method
$ \mathcal{Y} $ too large	exploit structure	belief propagation
$N$ too large	mini-batches	stochastic gradient descent
$K$ too large	trained $\psi$	piece-wise training

# Deep Structured Models

# Motivation

## Log-Linear Models:

$$p(y|x, w) = \frac{1}{Z(x, w)} \exp \{ \langle w, \psi(x, y) \rangle \}$$

- ▶ Log-linear in the parameters  $w \Rightarrow$  features must do all the heavy lifting
- ▶ Only linear combination of features is learned

## Deep Structured Models:

$$p(y|x, w) = \frac{1}{Z(x, w)} \exp \{ \psi(x, y, w) \}$$

- ▶ Potential functions directly parametrized via  $w$
- ▶ Results in a much more flexible model ( $\psi$  can represent, e.g., a neural network)

# Deep Structured Models

## Negative Log-Likelihood and its Gradient:

$$\begin{aligned}\mathcal{L}(w) &= -\sum_{n=1}^N \left[ \psi(x^n, y^n, w) - \log \sum_{y \in \mathcal{Y}} \exp \{ \psi(x^n, y, w) \} \right] \\ \nabla_w \mathcal{L}(w) &= -\sum_{n=1}^N \left[ \nabla_w \psi(x^n, y^n, w) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \nabla_w \psi(x^n, y, w) \right]\end{aligned}$$

- ▶ Similar form as for log-linear models
- ▶ Differences to log-linear model highlighted in red

# Deep Structured Models

## Negative Log-Likelihood and its Gradient:

$$\begin{aligned}\mathcal{L}(w) &= -\sum_{n=1}^N \left[ \psi(x^n, y^n, w) - \log \sum_{y \in \mathcal{Y}} \exp \{ \psi(x^n, y, w) \} \right] \\ \nabla_w \mathcal{L}(w) &= -\sum_{n=1}^N \left[ \nabla_w \psi(x^n, y^n, w) - \sum_{y \in \mathcal{Y}} p(y|x^n, w) \nabla_w \psi(x^n, y, w) \right]\end{aligned}$$

- Again, sums can be efficiently computed as features decompose

$$\psi(x, y, w) = \sum_{f \in \mathcal{F}} \psi_f(x, y_f, w)$$

- Let us now represent  $\psi_f(x, y_f, w)$  using deep neural networks



# Deep Structured Models

Algorithm:

- ▶ Forward pass to compute  $\psi_f(x, y_f, w)$
- ▶ Backward pass to obtain gradients  $\nabla_w \psi(x^n, y, w)$
- ▶ Compute marginals using message passing
- ▶ Update parameters  $w$

What is the problem with this approach?

- ▶ Very slow as inference in GM is required for every gradient update

Alternatives

- ▶ Interleave learning and inference [Chen *et al.*, ICML 2015]
- ▶ But still only applicable to very small scale problems

# Inference Unrolling

# Inference Unrolling

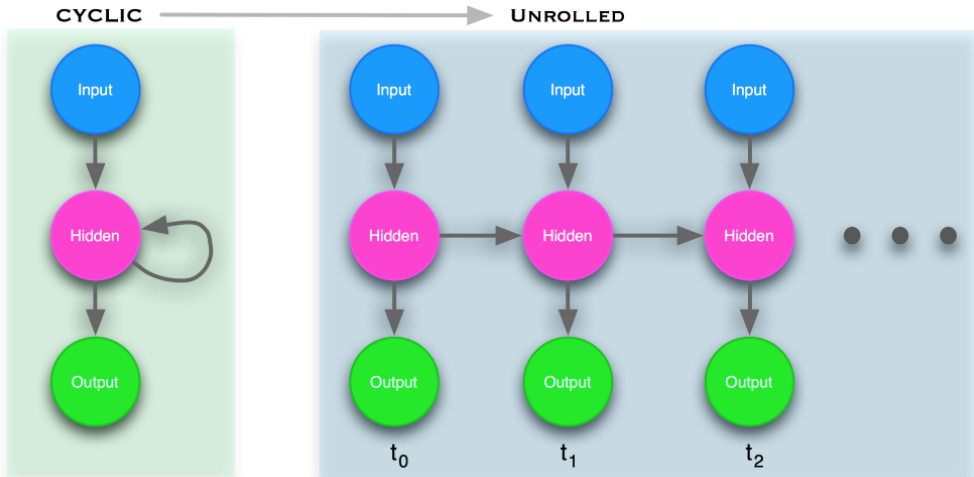
Idea:

- ▶ Consider inference as sequence of small computations
- ▶ “Unroll” a fixed number of inference iterations and consider as RNN
- ▶ Compute gradients, e.g., using automatic differentiation

Warning:

- ▶ Now: empirical risk minimization
- ▶ Thus purely deterministic approach, giving up probabilistic viewpoint
- ▶ But often fast enough for efficient training in deep models
- ▶ Effectively integrates structure of the problem into architecture of the network
- ▶ Can be thought of as a form of regularization (hard constraint)

# Inference Unrolling



# Automatic Differentiation

Idea:

- Rewrite complicated function as composition of simple functions:

$$f = f_0 \circ f_1 \circ \dots \circ f_n$$

- Each simple function  $f_k$  has a simple derivative


- Use chain rule:  $\frac{\partial f_0}{\partial f_1} \frac{\partial f_1}{\partial f_2} \dots \frac{\partial f_n}{\partial x}$

- Example:

$$f(x, y) = \cos(x) \sin(y) + \frac{x}{y}$$

<http://www.columbia.edu/~ahd2125/post/2015/12/5/>

Computation Graph:



Examples


# Conditional Random Fields as Recurrent Neural Networks

[Zheng *et al.*, ICCV 2015]

# Conditional Random Fields as Recurrent Neural Networks

$$E(\mathbf{x}) = \sum_i \psi_u(x_i) + \sum_{i < j} \psi_p(x_i, x_j), \quad (1)$$

$$\psi_p(x_i, x_j) = \mu(x_i, x_j) \sum_{m=1}^M w^{(m)} k_G^{(m)}(\mathbf{f}_i, \mathbf{f}_j), \quad (2)$$




---

**Algorithm 1** Mean-field in dense CRFs [29], broken down to common CNN operations.

---

$Q_i(l) \leftarrow \frac{1}{Z_i} \exp(U_i(l))$  for all  $i$  ▷ Initialization

**while** not converged **do**

$\tilde{Q}_i^{(m)}(l) \leftarrow \sum_{j \neq i} k^{(m)}(\mathbf{f}_i, \mathbf{f}_j) Q_j(l)$  for all  $m$  ▷ Message Passing

$\check{Q}_i(l) \leftarrow \sum_m w^{(m)} \tilde{Q}_i^{(m)}(l)$  ▷ Weighting Filter Outputs

$\hat{Q}_i(l) \leftarrow \sum_{l' \in \mathcal{L}} \mu(l, l') \check{Q}_i(l')$  ▷ Compatibility Transform

$\check{\check{Q}}_i(l) \leftarrow U_i(l) - \hat{Q}_i(l)$  ▷ Adding Unary Potentials


$Q_i \leftarrow \frac{1}{Z_i} \exp(\check{\check{Q}}_i(l))$  ▷ Normalizing

**end while**

---



# Conditional Random Fields as Recurrent Neural Networks



# RayNet: Learning Volumetric 3D Reconstruction with Ray Potentials

[Paschalidou, Ulusoy, Schmitt, van Gool & Geiger, CVPR 2018]

# RayNet: Learning Volumetric 3D Reconstruction


Distribution over voxel occupancies:

$$p(\mathbf{o}) = \frac{1}{Z} \prod_{i \in \mathcal{X}} \underbrace{\varphi_i(o_i)}_{\text{unary}} \prod_{r \in \mathcal{R}} \underbrace{\psi_r(\mathbf{o}_r)}_{\text{ray}}$$

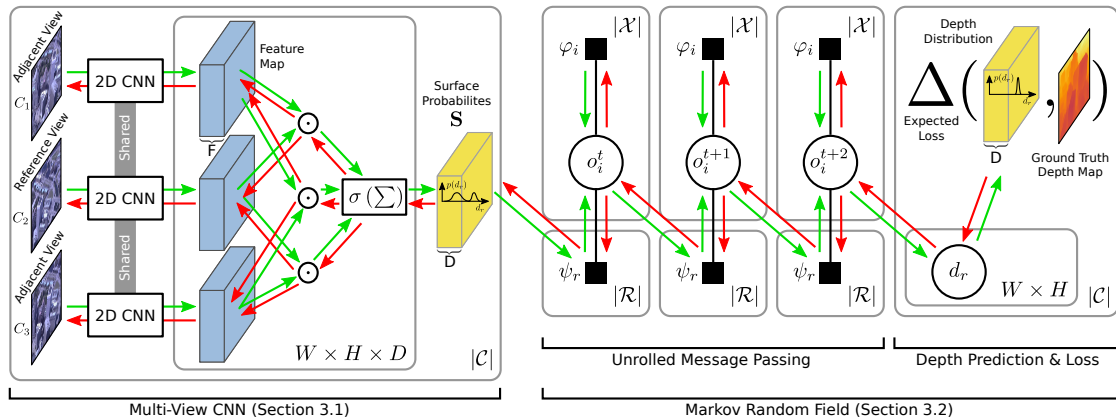
$$\varphi_i(o_i) = \gamma^{o_i} (1 - \gamma)^{1-o_i}$$

$$\psi_r(\mathbf{o}_r) = \sum_{i=1}^{N_r} o_i^r \prod_{j < i} (1 - o_j^r) s_i^r$$

Corresponding factor graph:




# RayNet: Learning Volumetric 3D Reconstruction



# RayNet: Learning Volumetric 3D Reconstruction



(a) Image




(b) Ours (CNN)




(c) Ours (CNN+MRF)



(d) ZNCC



(e) Ulusoy et al. [35]



(f) Hartmann et al. [14]

Q & A Session